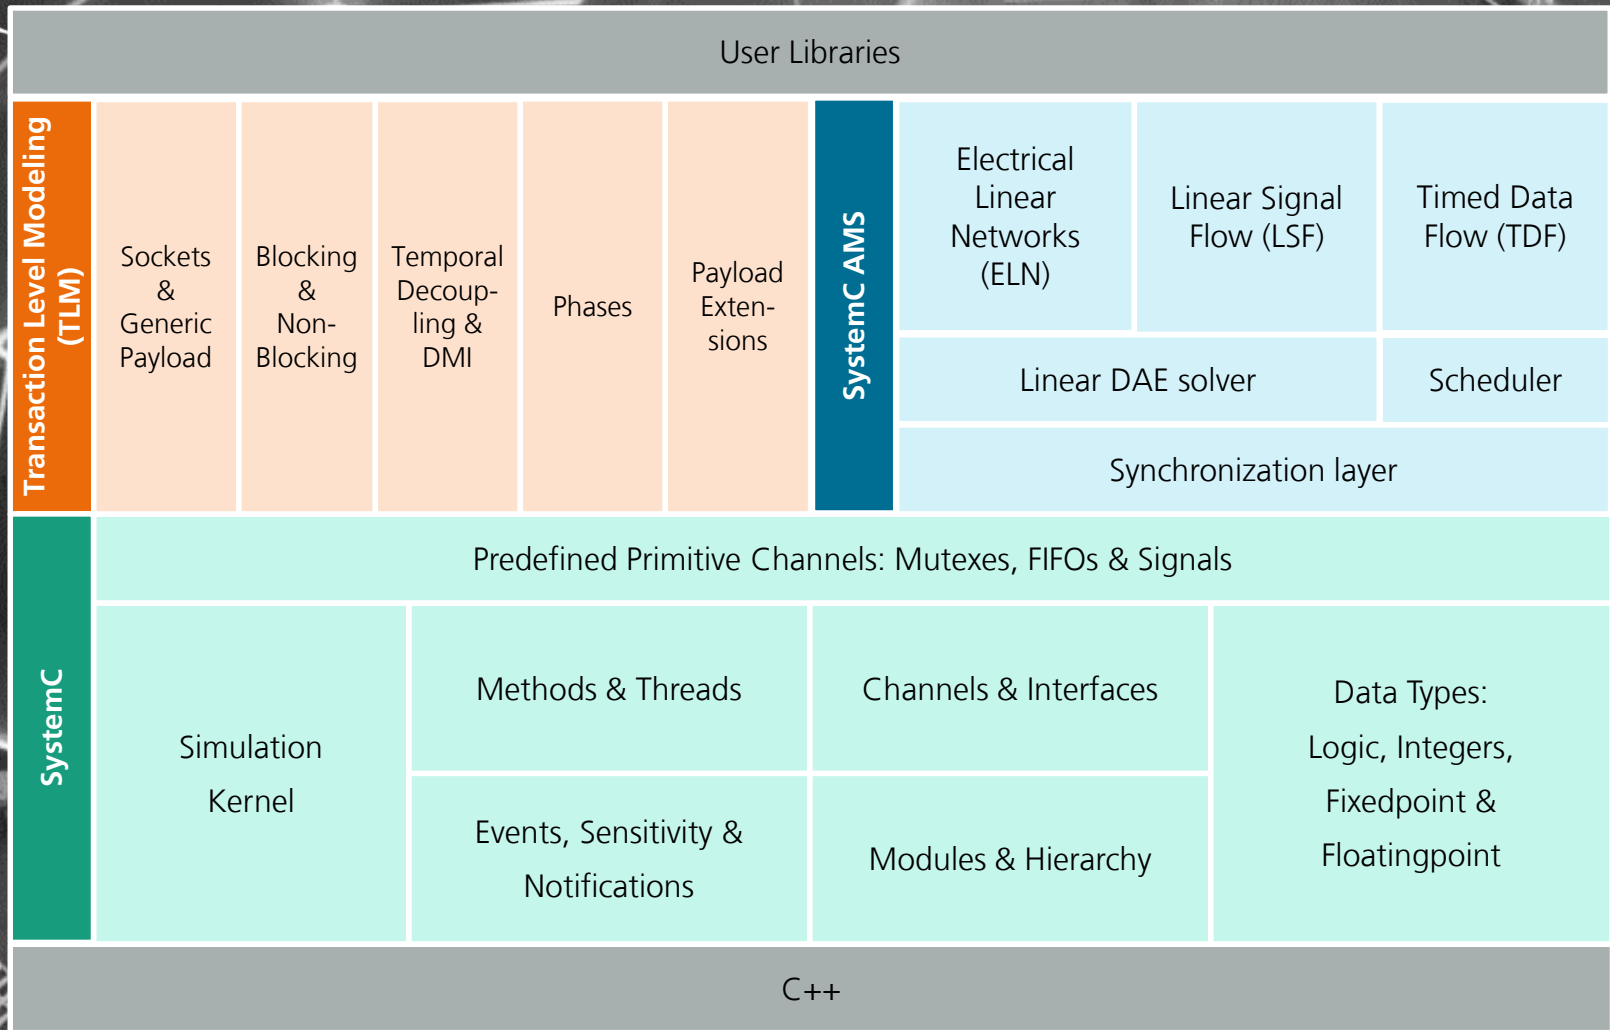


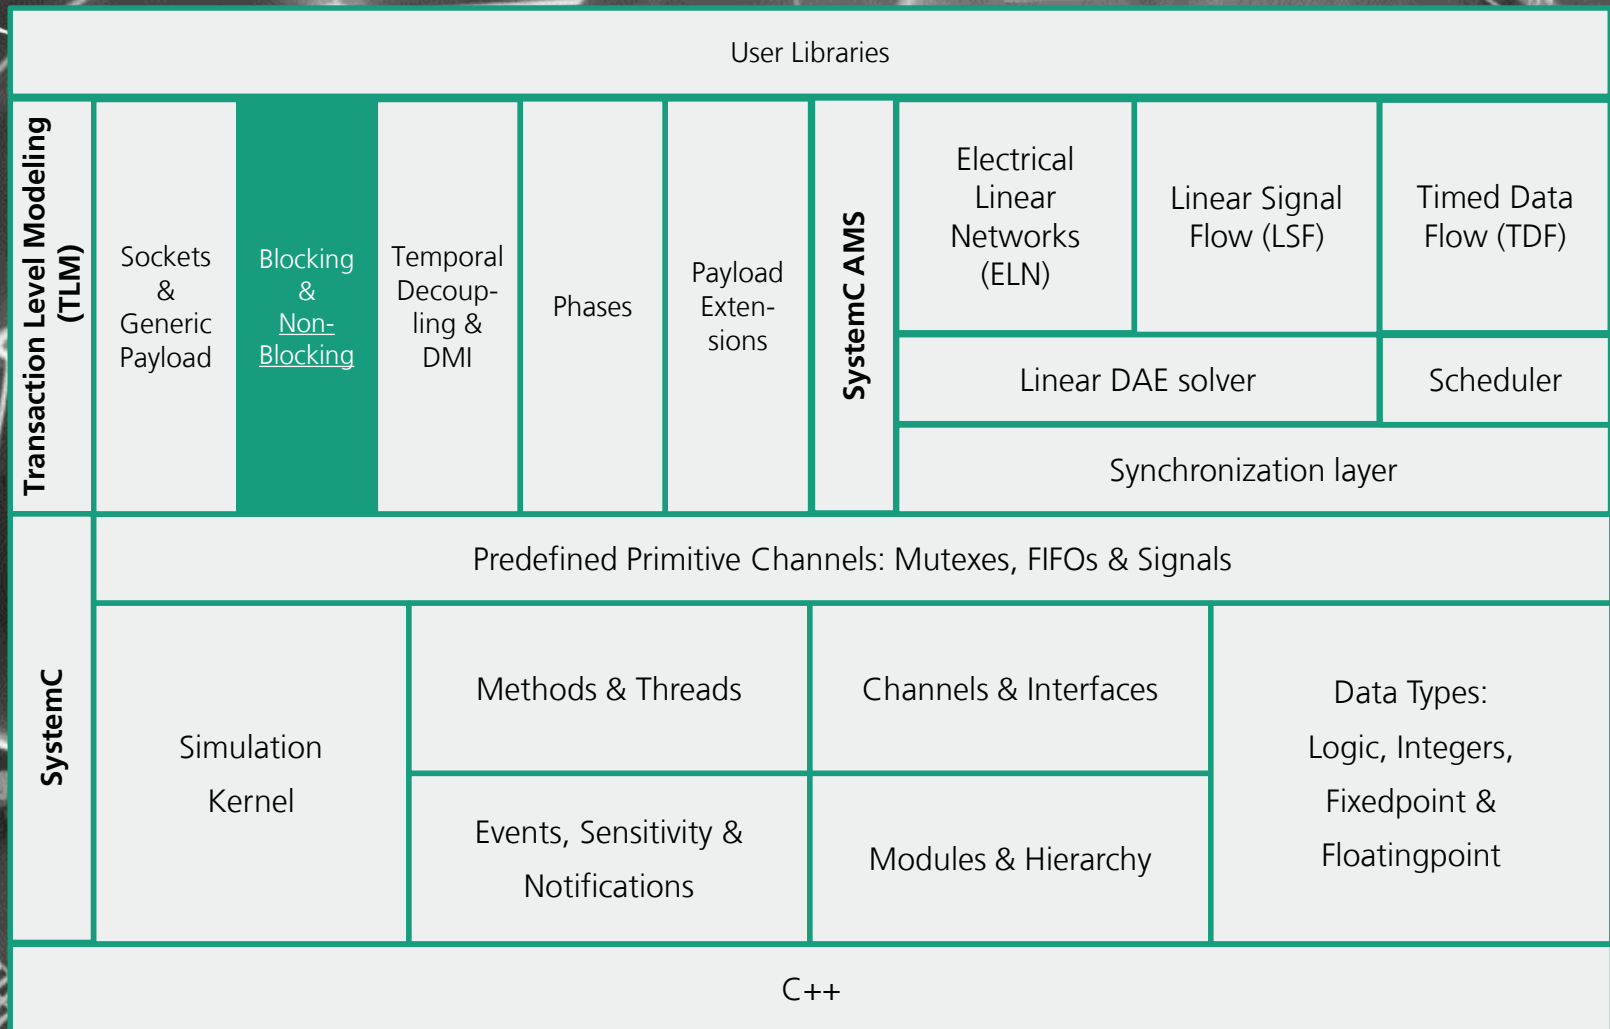
SystemC and Virtual Prototyping

Dr. Matthias Jung, Fraunhofer Institute IESE

matthias.jung@iese.fraunhofer.de







Recap: TLM Coding Styles and Mechanisms

TLM Use Cases

SW Application
Development

SW Performance
Analysis

Architecture
Analysis

Hardware
Verification

TLM 2.0 Coding Style *(Just Guidelines)*

Loosely-Timed (LT)

Single-phase, blocking API

Multi-phase, non-blocking API

Approximately-Timed (AT)

TLM Mechanisms *(Definitive API for enabling Interoperability)*

Blocking
transport

DMI

Quantum
(Keeper)

Sockets


Generic
payload

Extensions

Phases

Non-
blocking
transport

Recap: TLM Coding Styles and Mechanisms



Recap: Coding Styles in TLM

■ Loosely-Timed (LT):




- As fast as possible
- Sufficient timing detail to boot OS and run multicore systems and to develop SW or drivers
- Processes can run ahead of simulation time (temporal decoupling)
- Each transaction completes in one blocking function call
- Usage of Direct Memory Interface (DMI) e.g. for boot process

■ Approximately-Timed (AT):




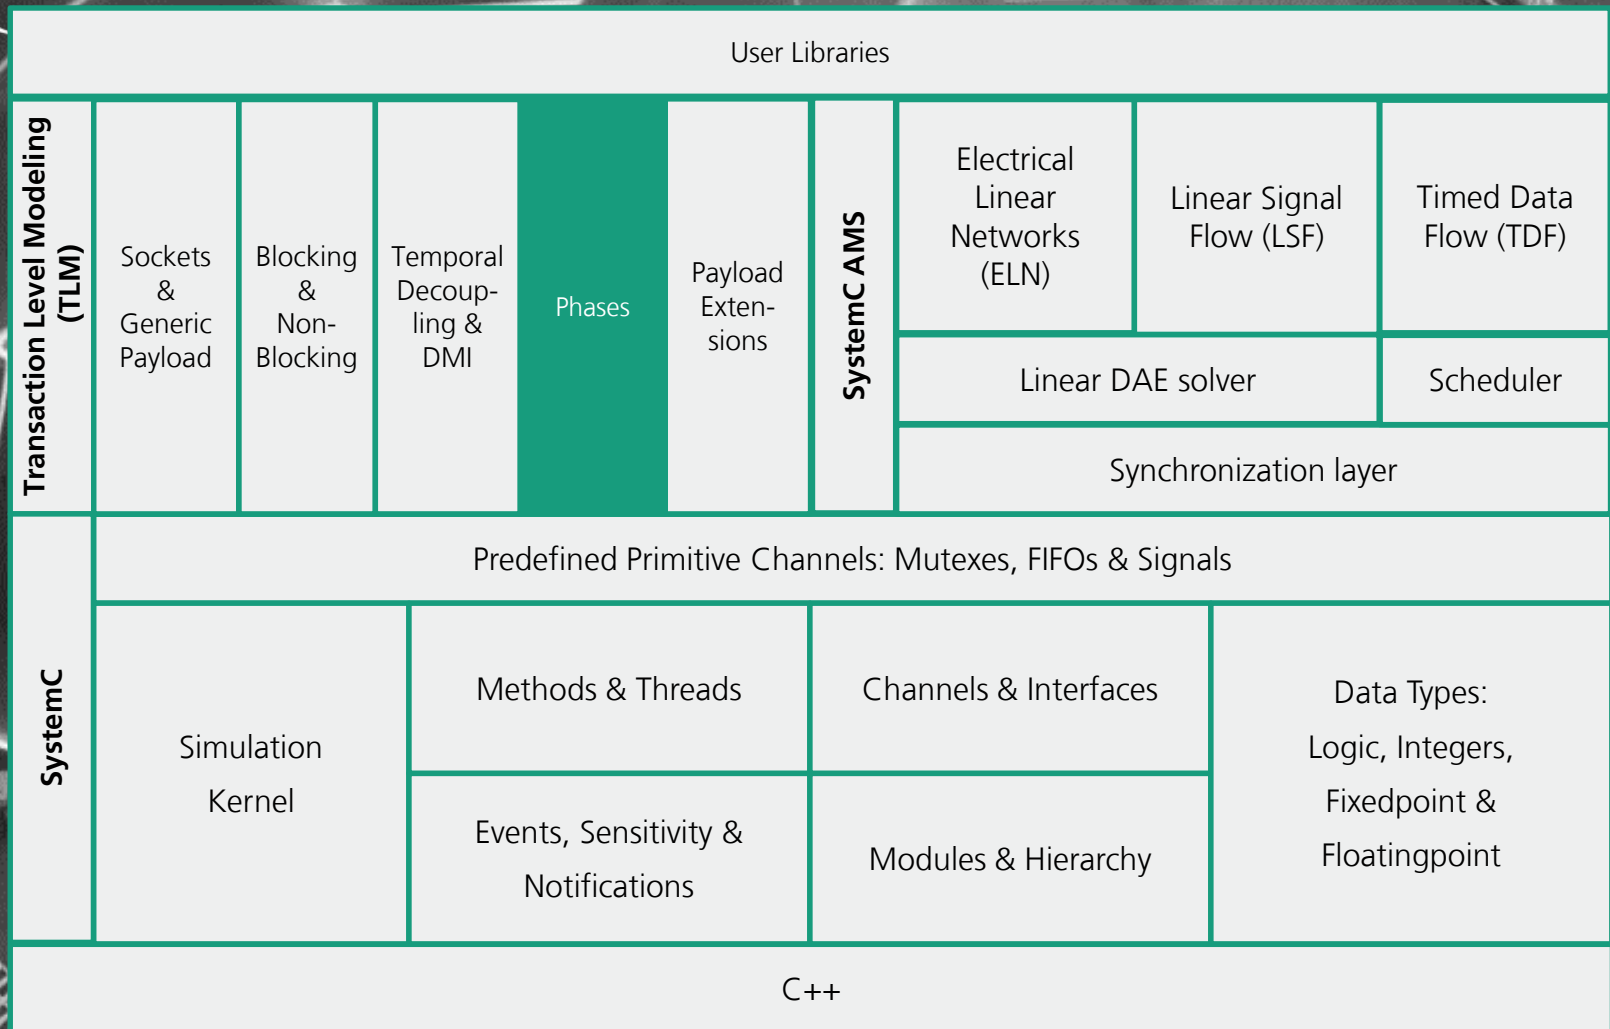
- Accurate enough for performance modelling
- Sufficient for architectural HW design space exploration
- Processes run in lockstep with simulation time
- Each transaction has usually 4 timing points i.e. 4 function calls (extensible if required, also less possible); non-blocking behavior
- More detailed than LT and therefore also slower than LT

Initiators, Targets and Interconnect

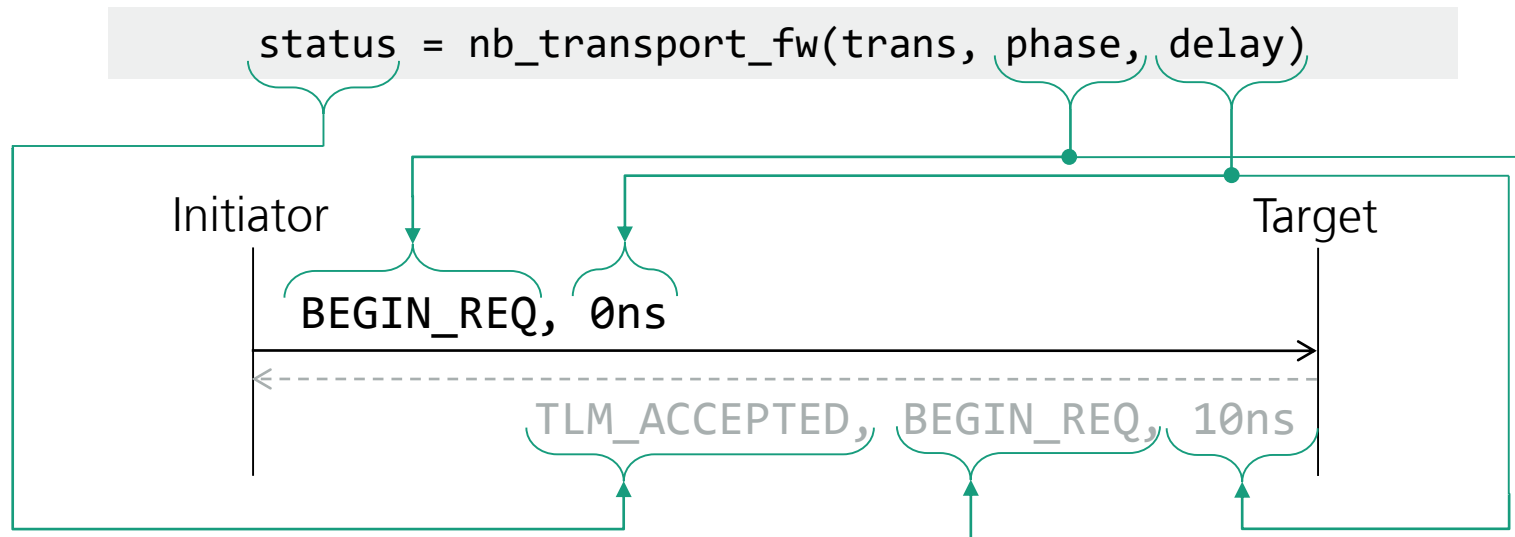


- References to the object are passed along the forward and backward paths:
 - LT uses Forward and Return Path
 - AT uses Forward, Backward, FW Return Path, BW Return Path
- AT uses non-blocking transport
- Time is handled with Payload Event Queues (PEQs)
- Allows modelling of O-O-O Cores and Backpressure
- Base Protocol

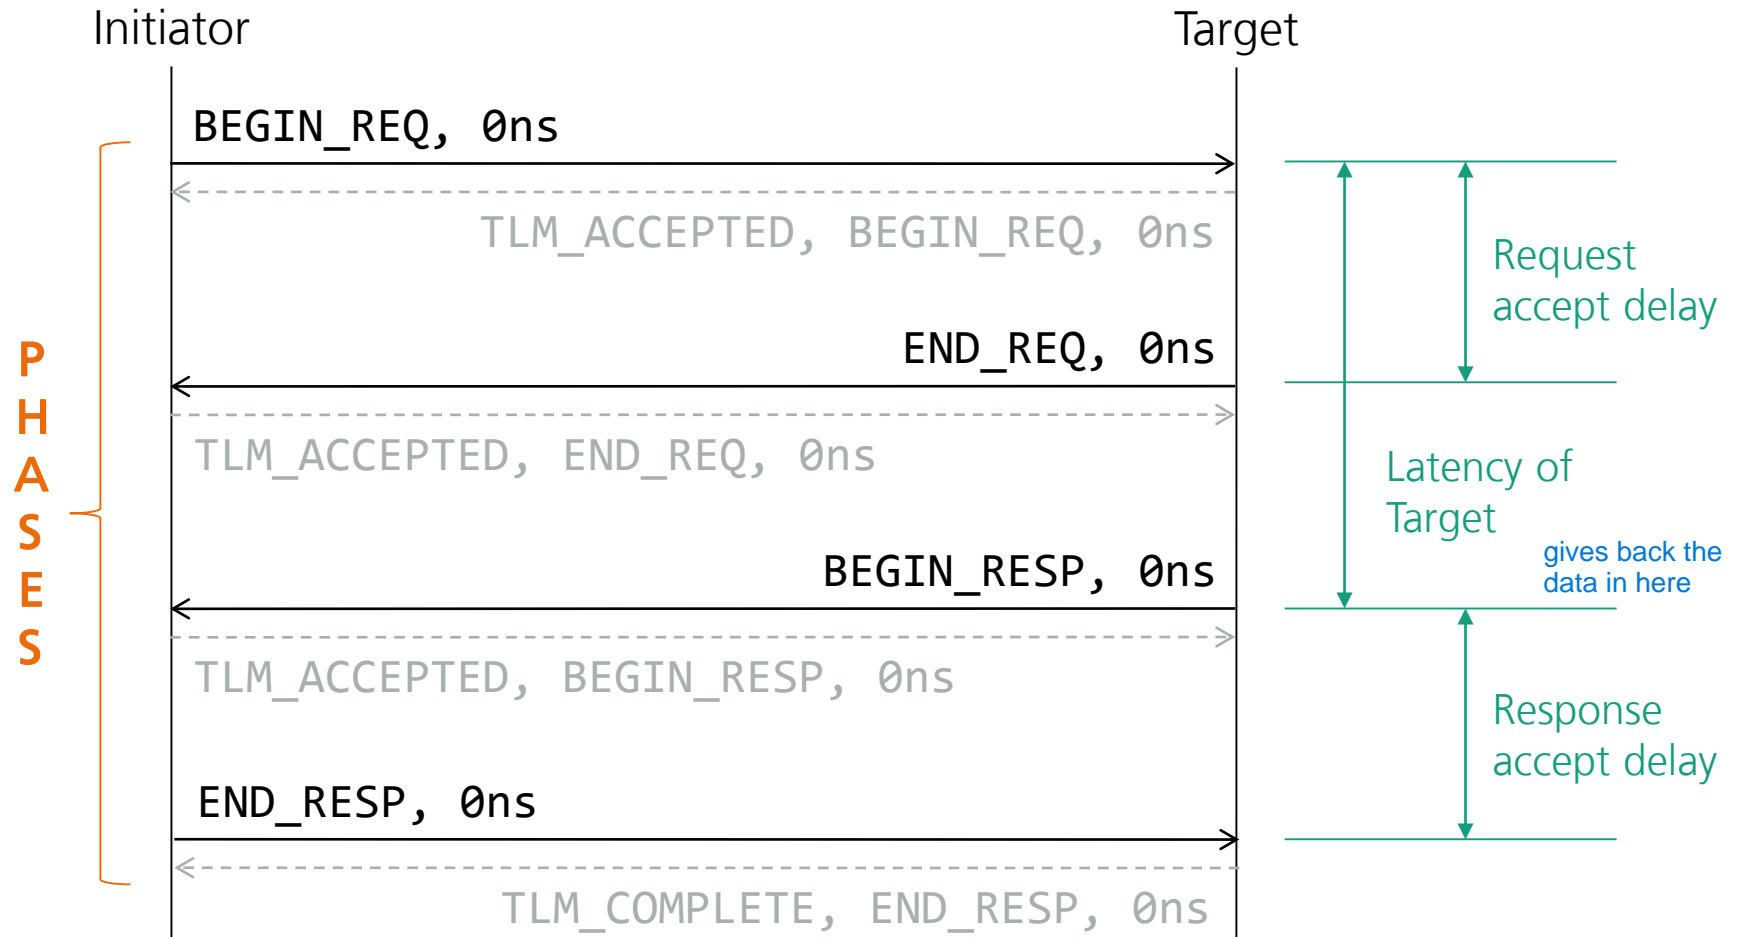




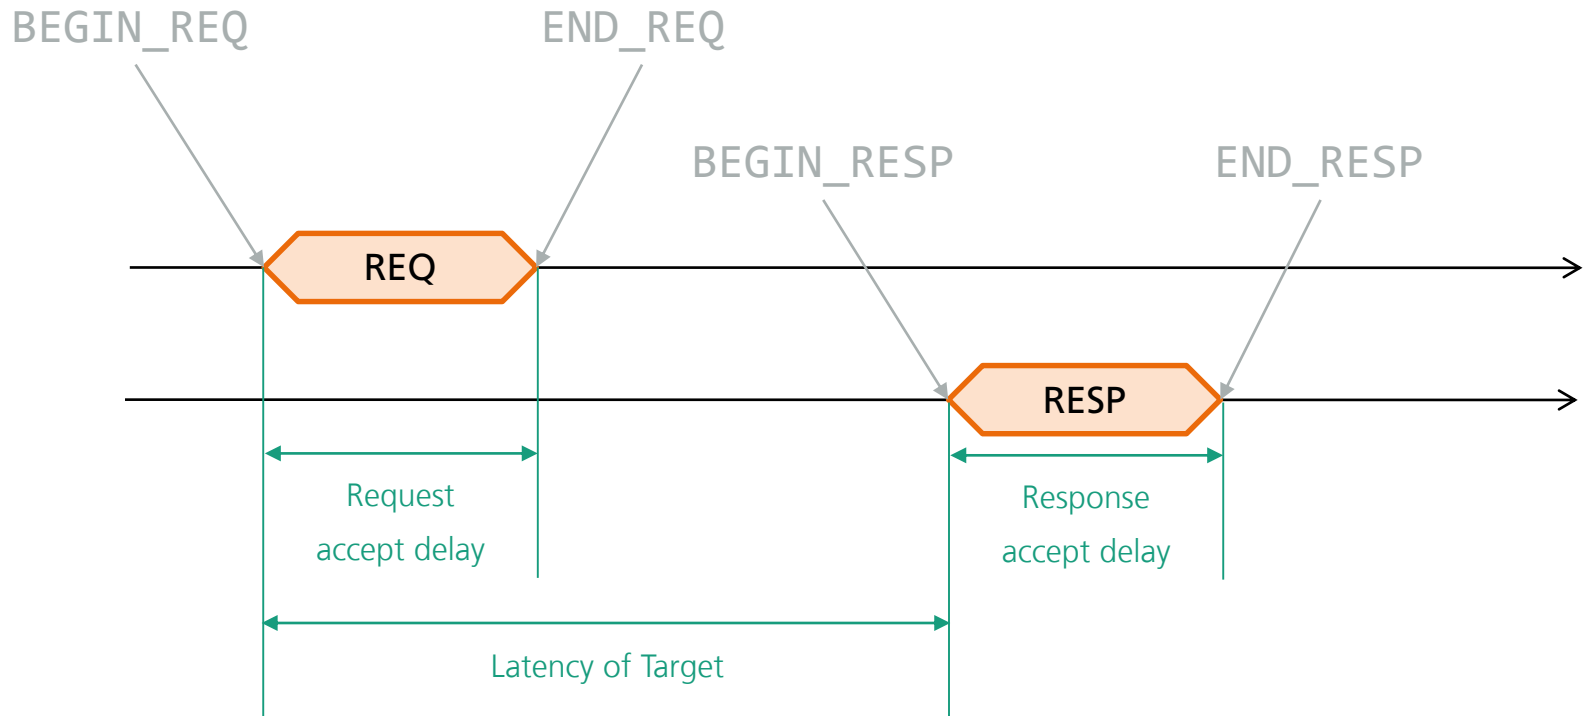
Non-Blocking Transport (AT) Base Protocol



Base Protocol Rules [1]: Using BW Path



Alternative View



Base Protocol Rules: Using BW Path


- Base Protocol phases
 - BEGIN_REQ → END_REQ → BEGIN_RESP → END_RESP
 - Must occur in increasing simulation time order
 - Phases must change with each call
- nb_transport_fw must not call nb_transport_bw directly and vice versa (PEQ!)
- Generic Payload memory management rules (See TLM Advanced)
- Extensions must be ignorable (See TLM Advanced)
- Target should handle mixed b_transport / nb_transport

Base Protocol Rules: The Exclusion Rule



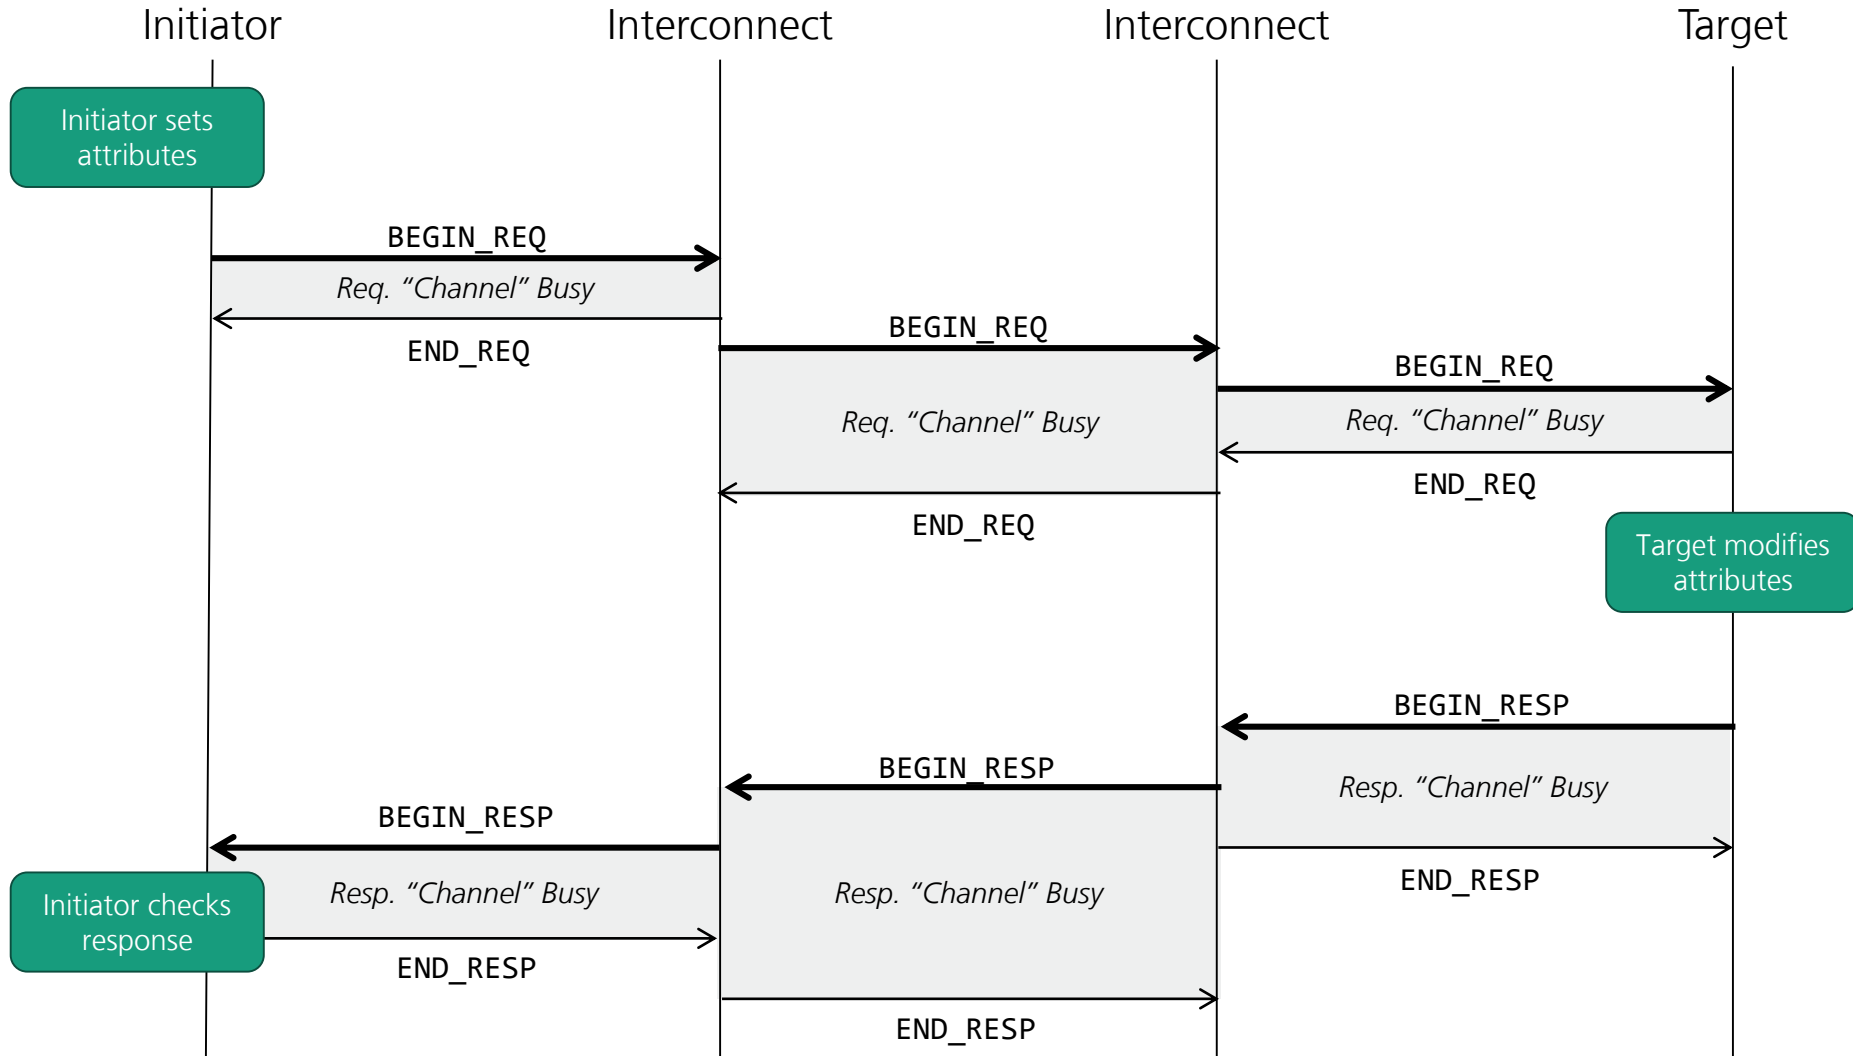
- Initiators and targets must honor the *Exclusion Rule*:
 - An Initiator must not send a new request (**BEGIN_REQ**) until it has received the **END_REQ** from the previous transaction
 - An target must not send the next **BEGIN_RESP** until it has received the **END_RESP** from the previous transactions
- The exclusion rules enable flow control like back pressure:
 - E.g. if an input buffer of a target is full the target can defer the sending of **END_REQ** for the transaction that filled the buffer, until the buffer has available space again
 - An RTL ready signal can be modeled by deferring **END_REQ**
 - However, since it is non-blocking, the target can do something else, e.g. sending

Modelling of Backpressure



- The exclusion rules enable flow control like back pressure:
 - E.g. if an input buffer of a target is full the target can defer the sending of **END_REQ** for the transaction that filled the buffer, until the buffer has available space again
 - An RTL ready signal can be modeled by deferring **END_REQ**
 - However, since it is non-blocking, the target can do something else, e.g. sending
- Also Response exclusion rule must be honored!

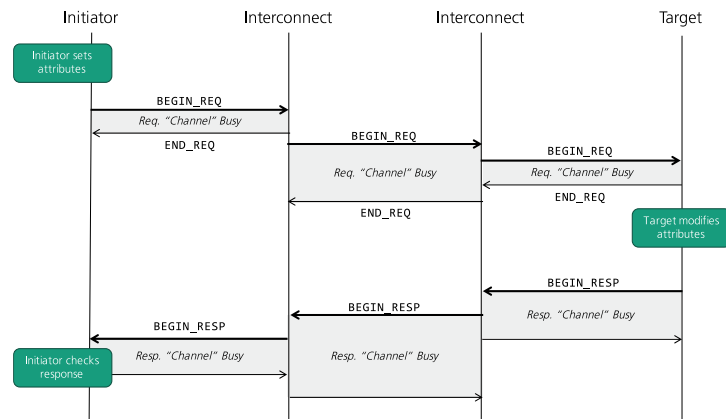
Base Protocol Rules: Causality with nb_transport



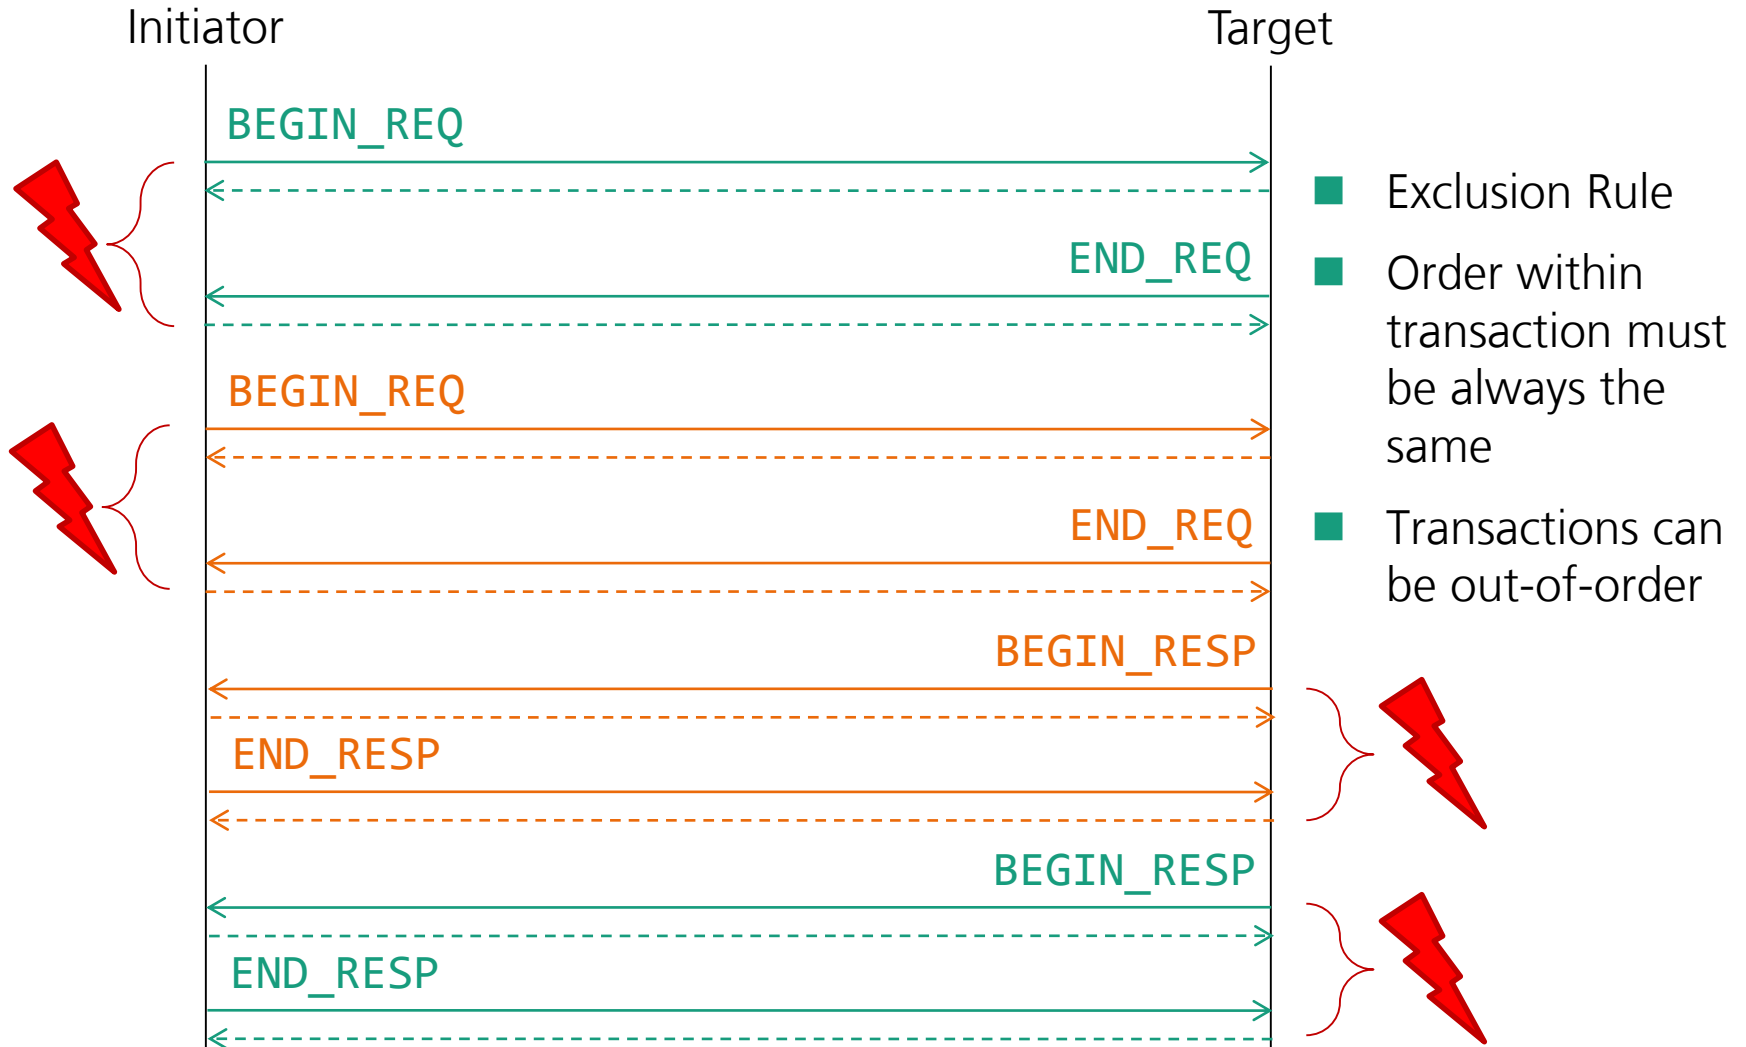
Base Protocol Rules: Causality with nb_transport




- BEGIN_REQ and BEGIN_RESP are propagated from end-to-end
- END_REQ and END_RESP are not propagated → they are local to each hop
- The END_REQ and END_RESP phases indicate that the requests and response “channels” are no longer busy and may be used for the next request or response respectively
- Initiator should not check the transaction payload until it has received the response

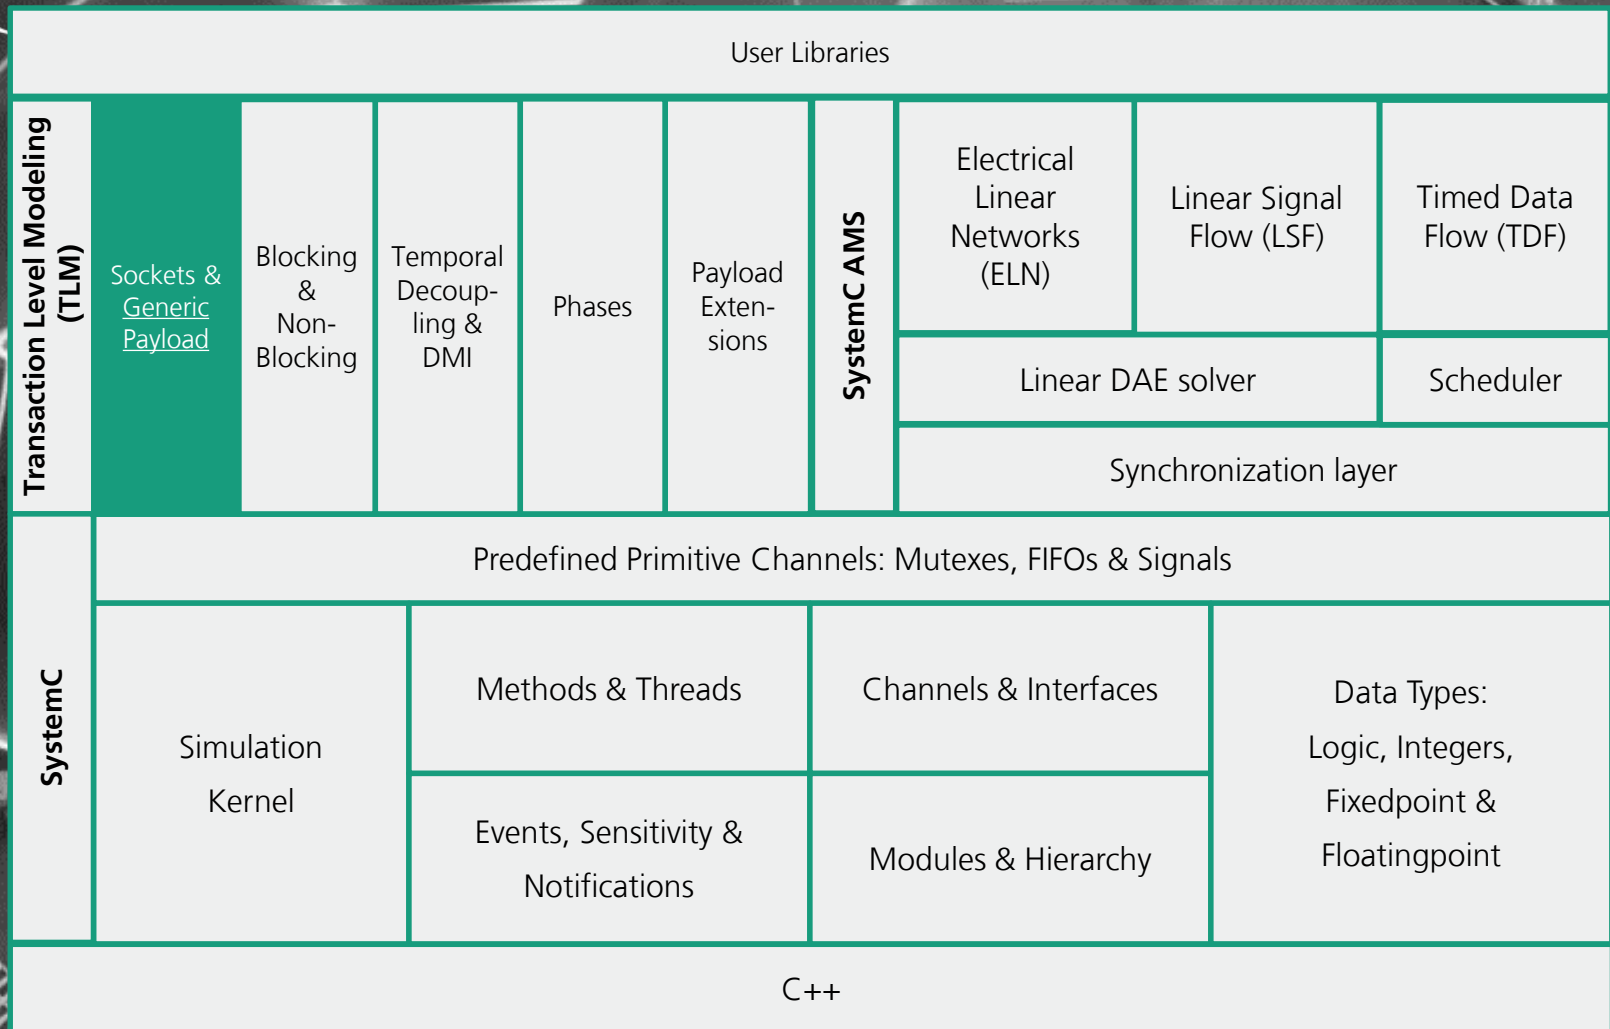


Base Protocol Rules: Pipelining of Transactions



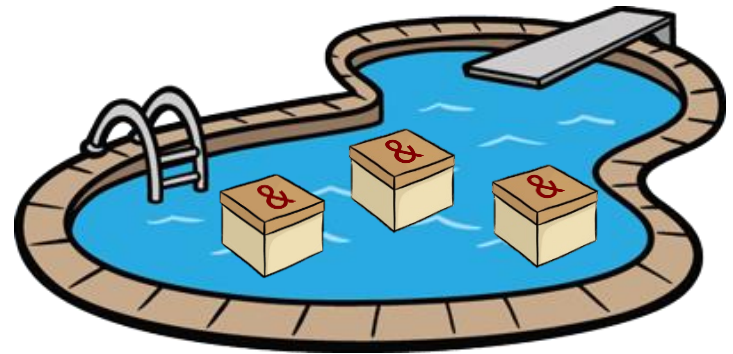
Alternative View





Generic Payload Pools (a.k.a Memory Management)

- Allocating of generic payload objects consumes wall-clock time
- Idea: do not allocate for each transaction a new generic payload → **Reuse**
 - A *Memory Manager* handles the generic payload objects in a pool
 - Before sending a transaction a payload object is **allocated** from the pool
 - Modules that send (or receive) this payload object must increase a reference count (**acquire**), which signalizes the memory manager that this object is still in use.
 - If a module is finished with the payload object, the reference count is decreased (**release**)
 - If the reference count is 0 the payload is freed and goes back into the pool.
 - If all transactions in the pool are in use a new one is generated



Memory Manager in TLM

- The memory manager is not part of the SystemC TLM Standard
- The TLM Standard provides only an interface class:

```
class tlm_mm_interface
{
    public:
    virtual void free(tlm::tlm_generic_payload*) = 0;
    virtual ~tlm_mm_interface(){}
}
```

- The implementation of the memory manager is up to the end user
- Virtual Platform tools like Synopsys Platform Architect have clever implementations

Example Memory Manager

Use code on github:

https://github.com/TUK-SCVP/SCVP.artifacts/tree/master/tlm_memory_manager

```
class MemoryManager : public tlm::tlm_mm_interface {
private:
    unsigned int numberOfAllocations;
    unsigned int numberOfFrees;
    std::vector<tlm::tlm_generic_payload*> freePayloads;

public:
    MemoryManager(): numberOfAllocations(0), numberOfFrees(0) {}

    ~MemoryManager(){
        for(tlm::tlm_generic_payload* payload: freePayloads) {
            delete payload;
            numberOfFrees++;
        }
    }

    tlm::tlm_generic_payload* MemoryManager::allocate(){
        if(freePayloads.empty()) {
            numberOfAllocations++;
            return new tlm::tlm_generic_payload(this);
        } else {
            tlm::tlm_generic_payload* result = freePayloads.back();
            freePayloads.pop_back();
            return result;
        }
    }

    void free(tlm::tlm_generic_payload* payload) {
        payload->reset(); //clears all fields and extensions
        freePayloads.push_back(payload);
    }
};
```

Transaction Pool
implemented as
`std::vector`

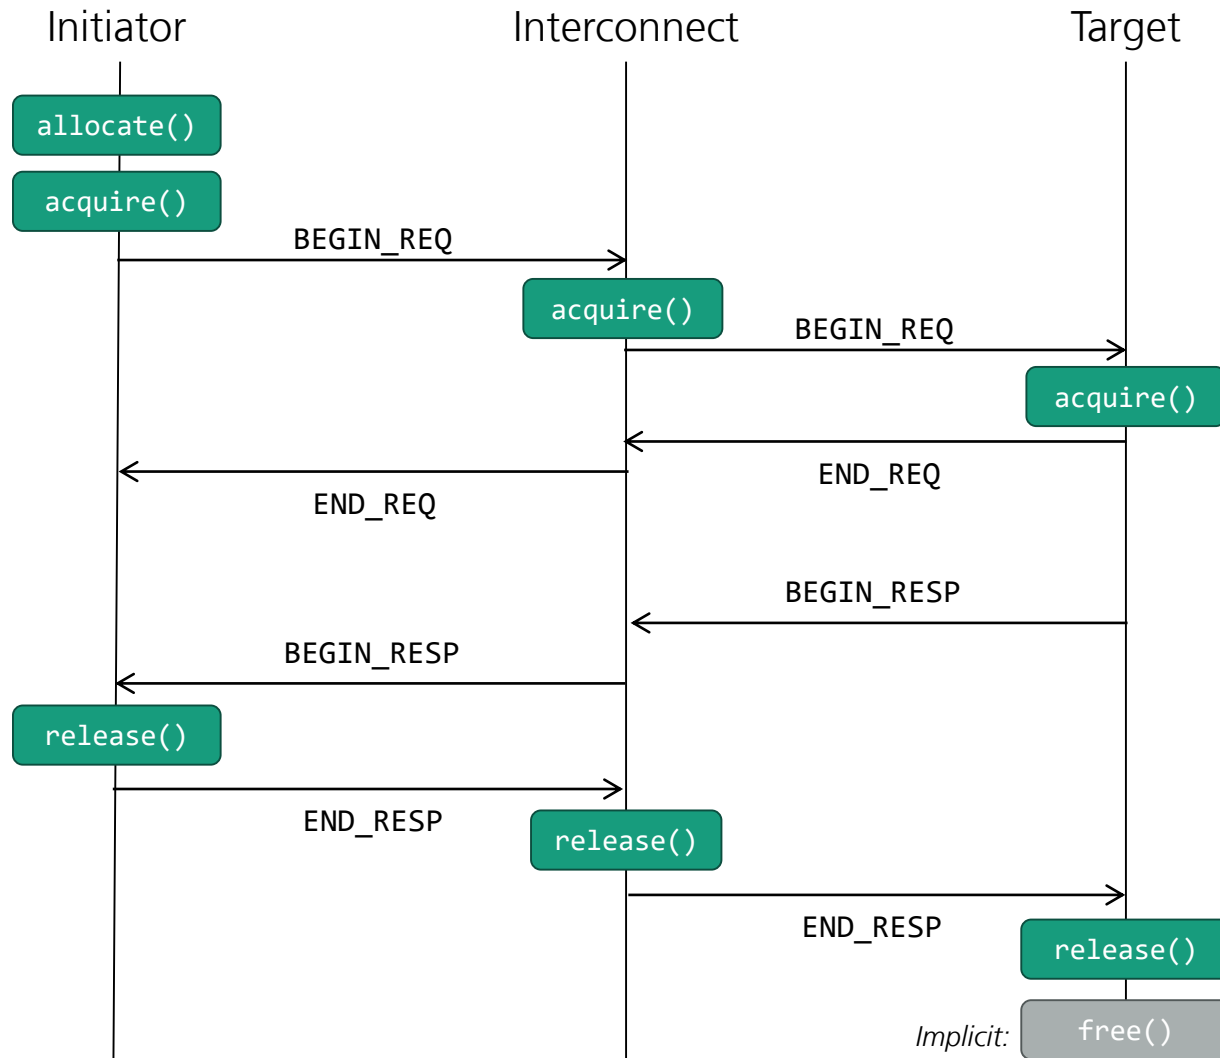
Cleanup
transaction pool

Allocate new
payload

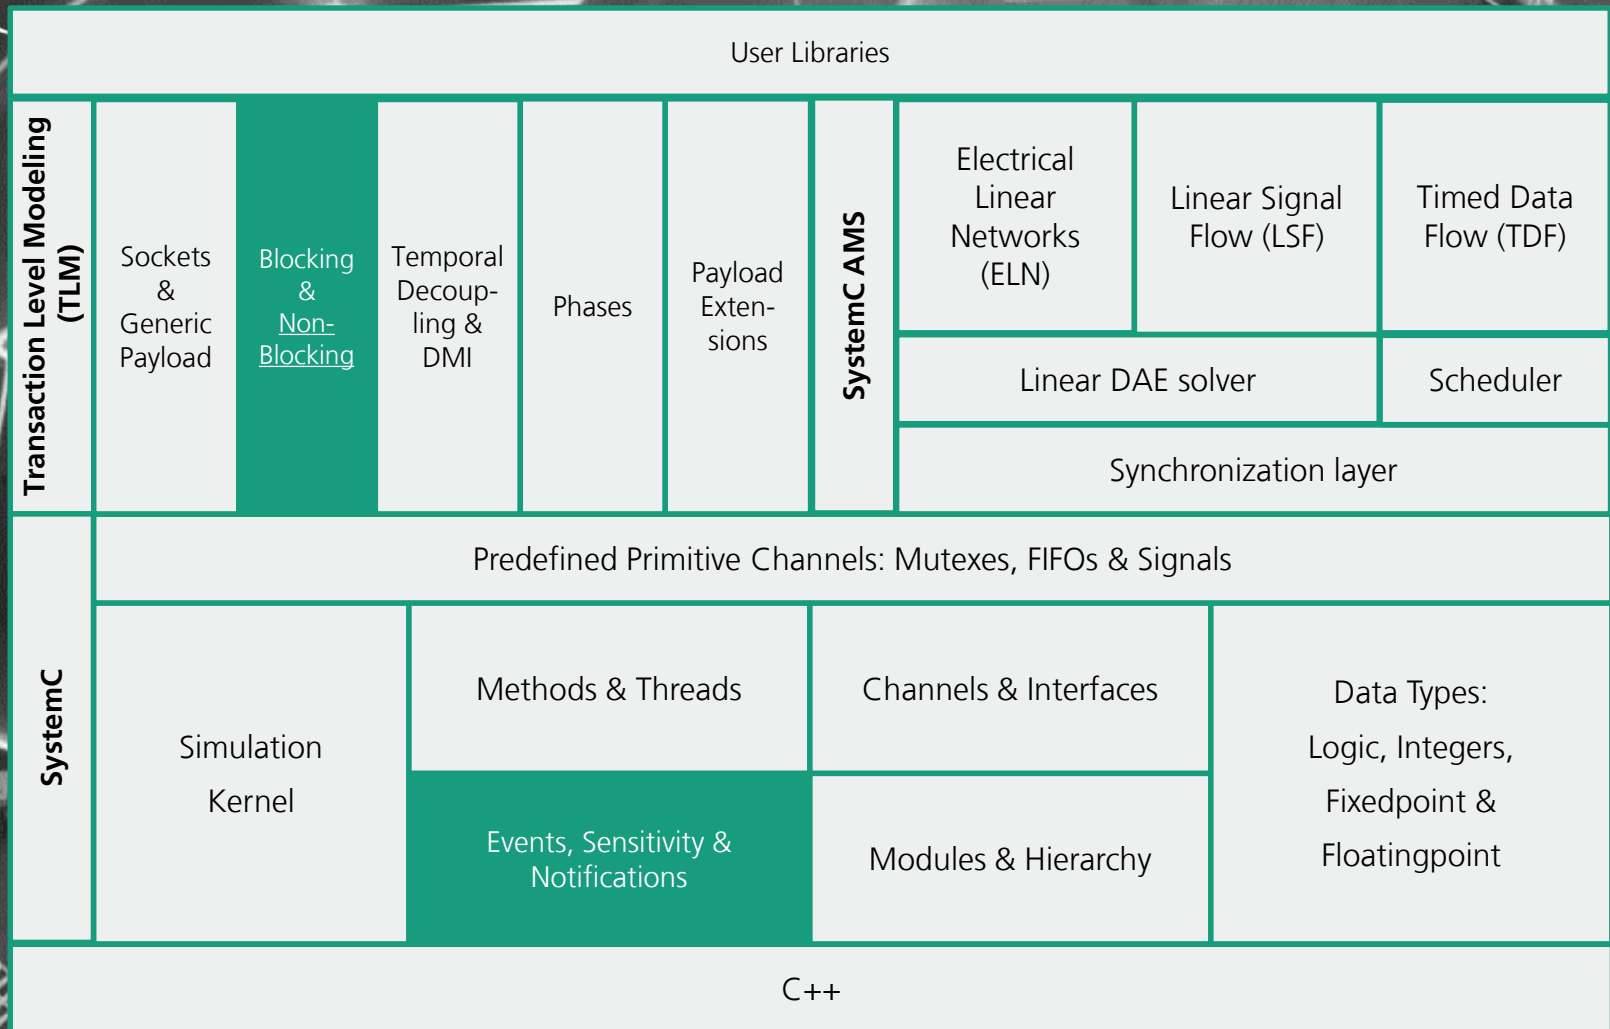
Reuse Payload
and remove
from pool

Add transaction
back to the pool

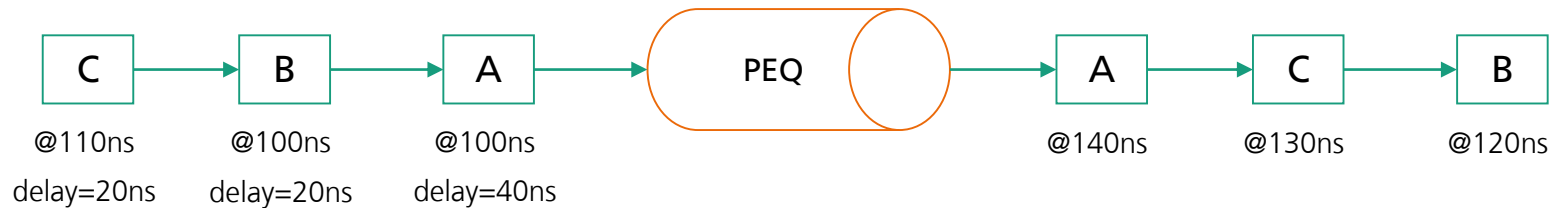
Usage of Memory Manager



- **b_transport**
a.k.a. LT may use a memory manager
- **nb_transport**
a.k.a. AT must use a memory manager

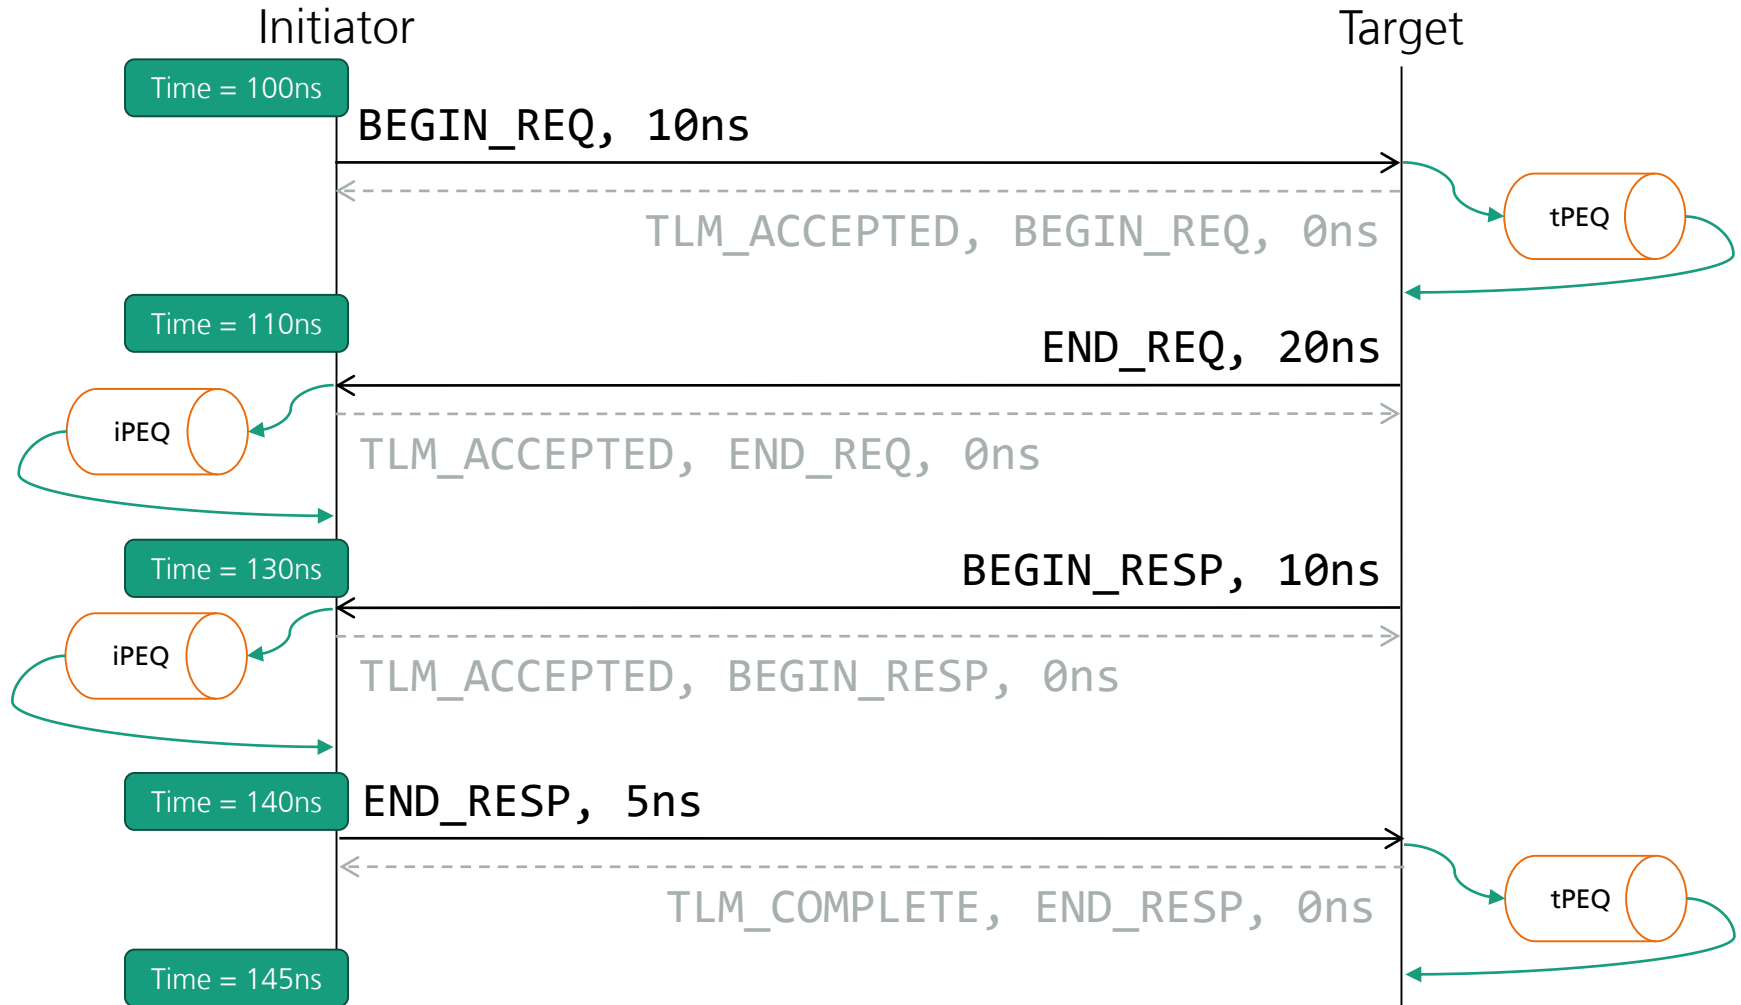


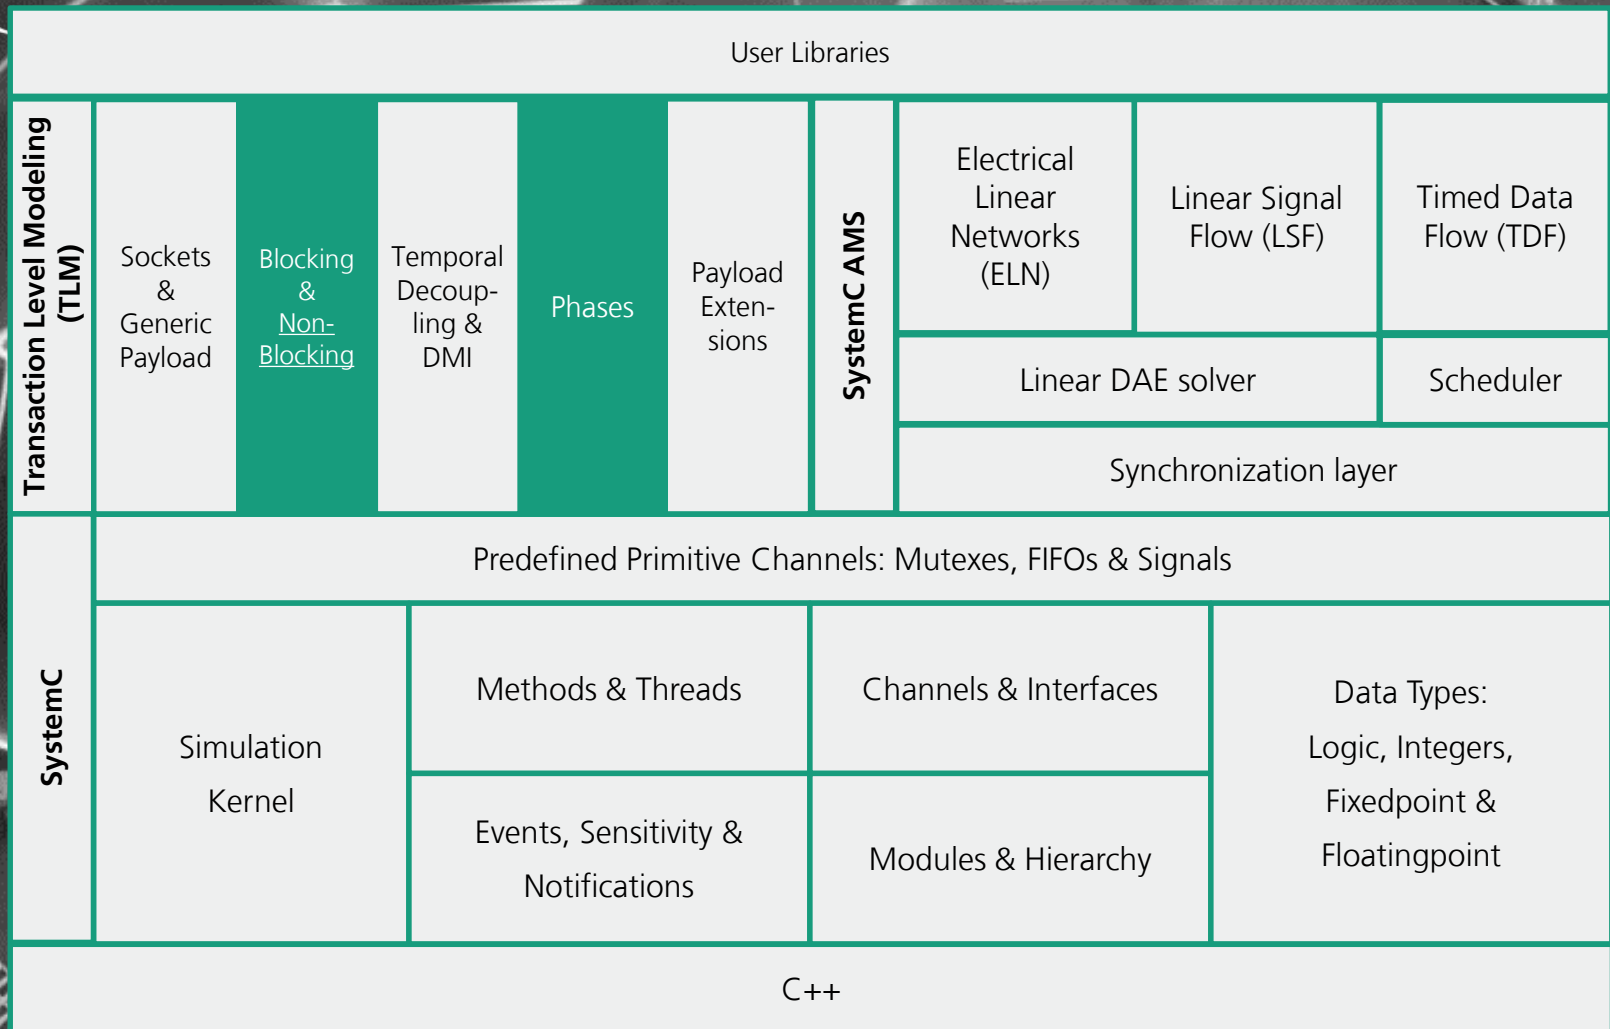
The Payload Event Queue



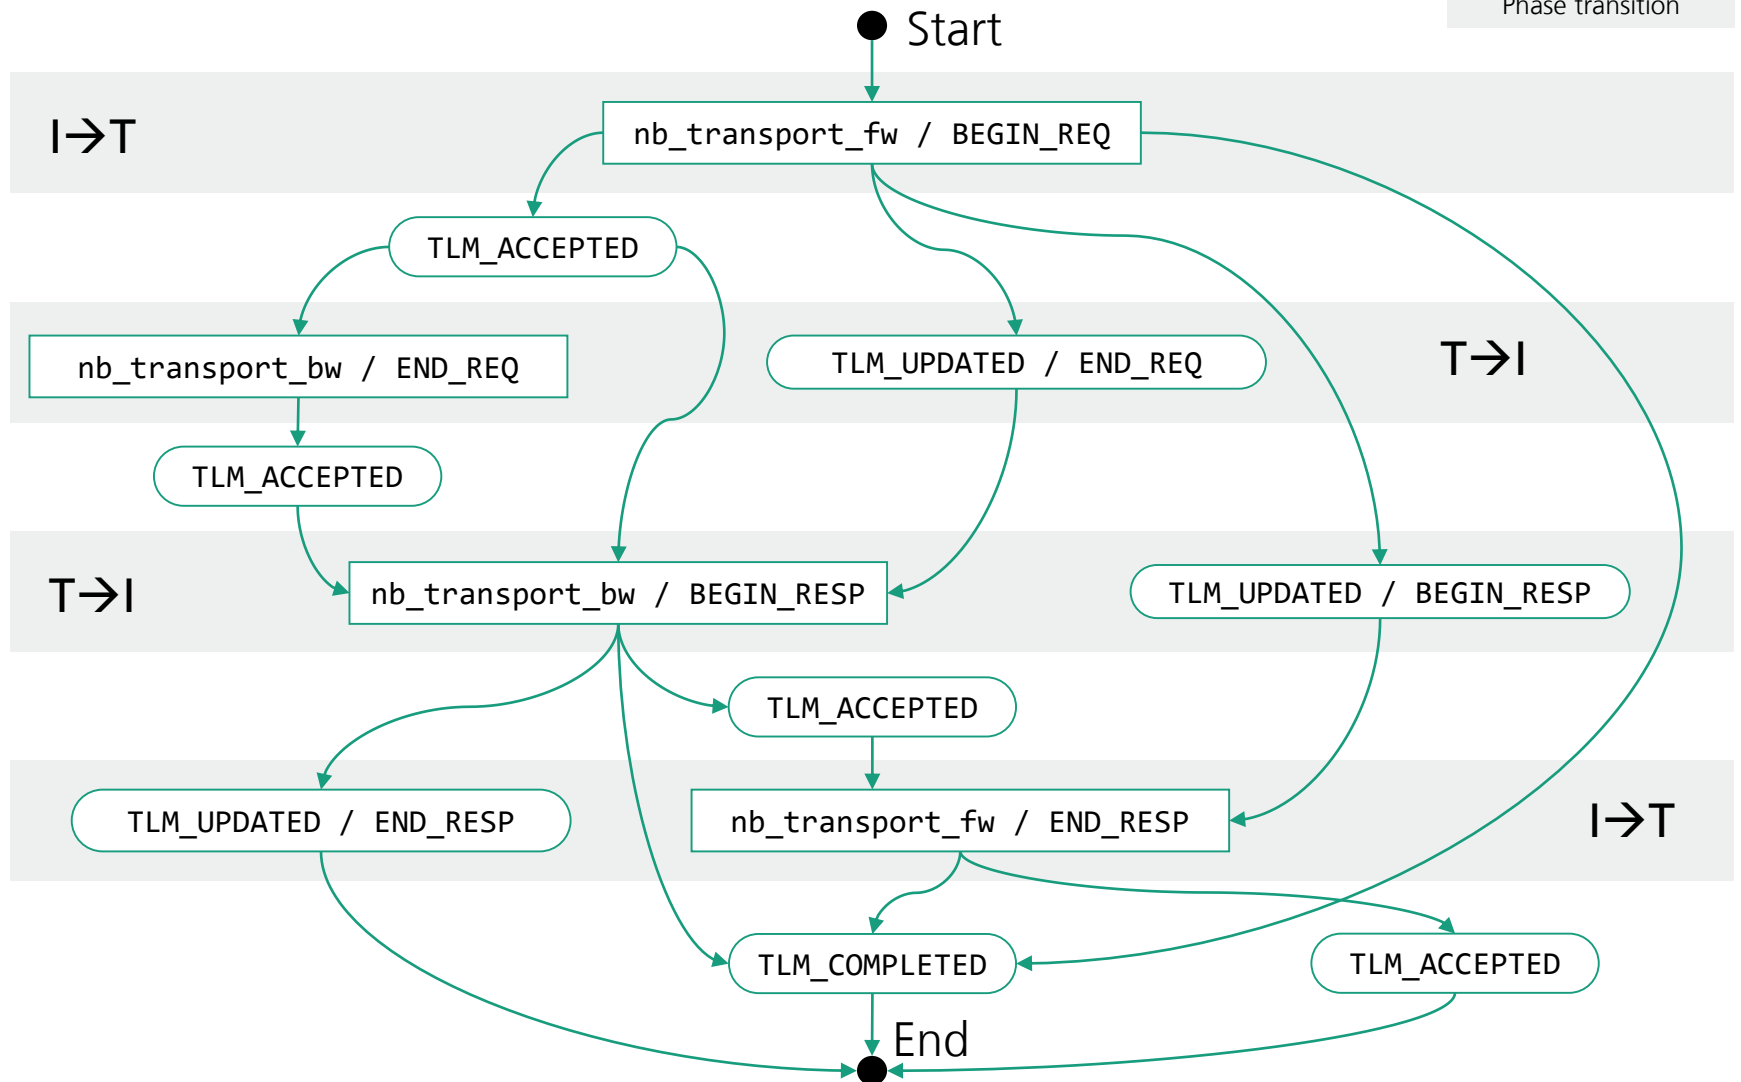
- AT models usually synchronize incoming calls with the simulation time
- Timing annotation is similar to b_transport:
 - *Hey target, please pretend that you received this message 10ns in the future*
- An AT component is usually not calling `wait()` statements for synchronizing with time, it rather posts the incoming transaction into a *Payload Event Queue* (PEQ). The PEQ internally synchronizes with the simulation time and calls a callback function in order to process the transaction when the time is reached.
- A component receives a transaction that should be processed in the future, so it puts the transaction into a PEQ in order to process it when time is reached

Timing Annotation with AT Models





Full Overview of AT Protocol Possibilities

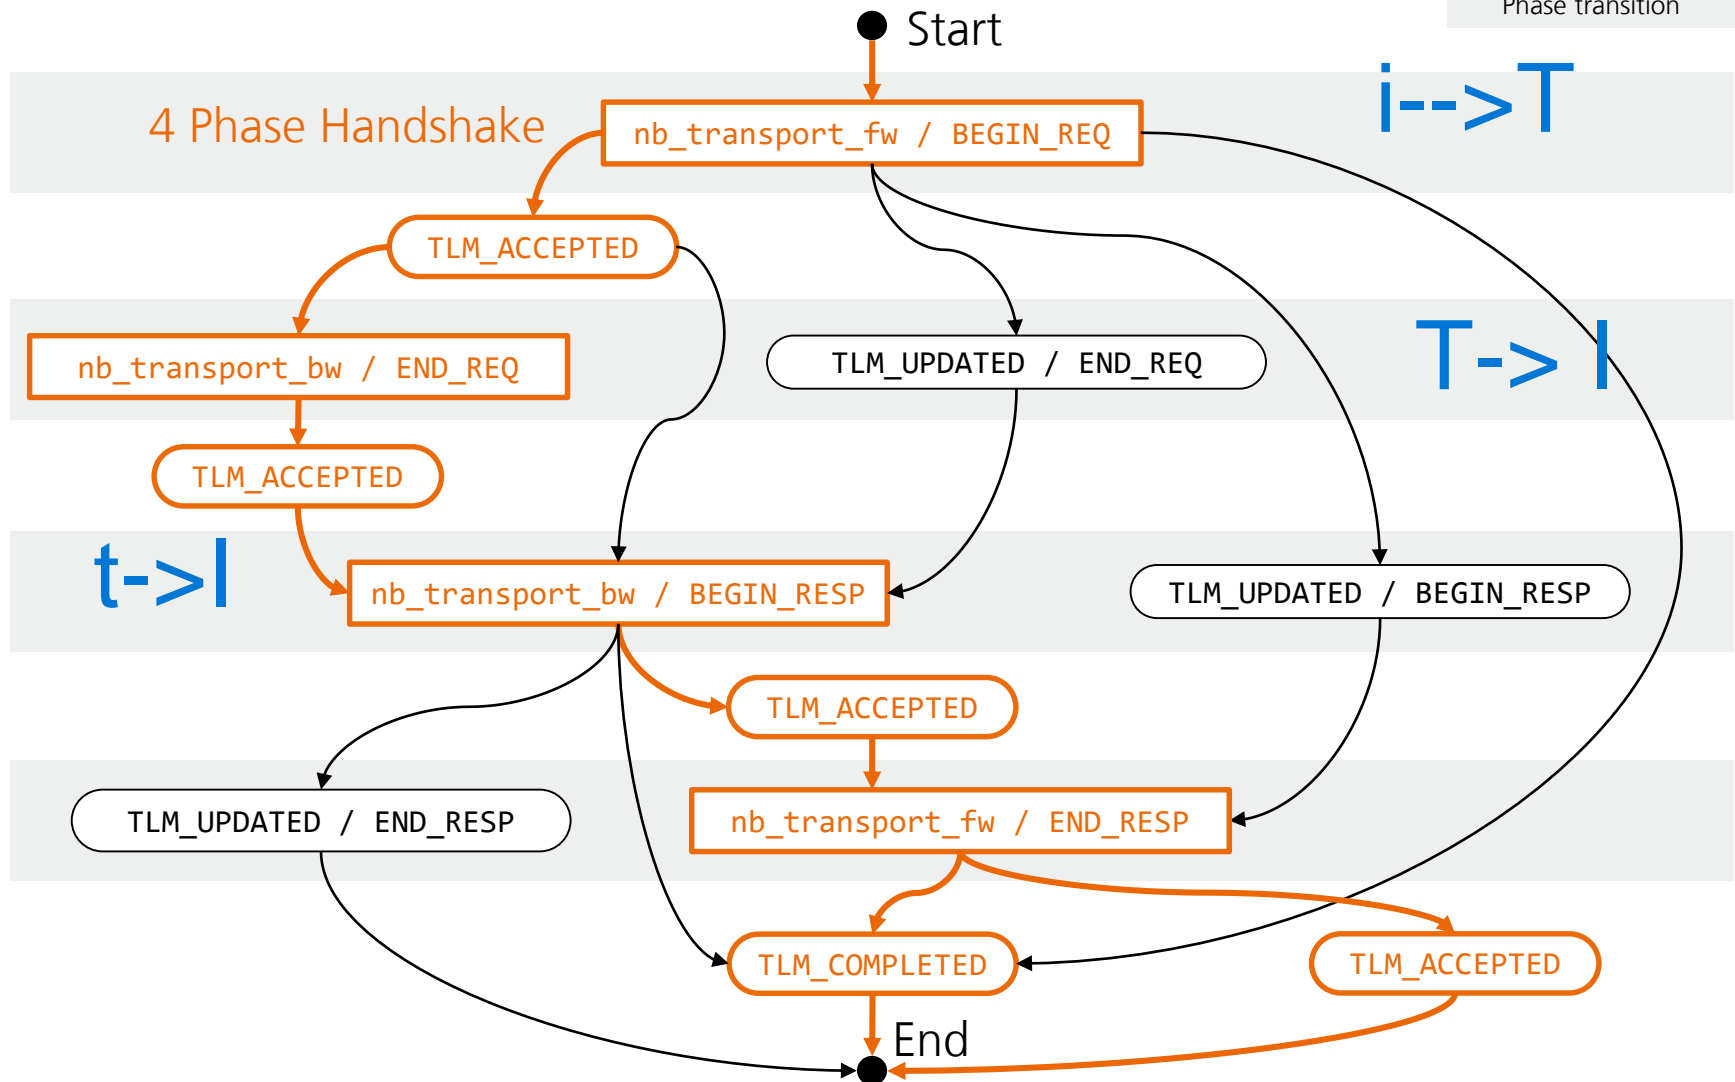


Full Overview of AT Protocol Possibilities

call

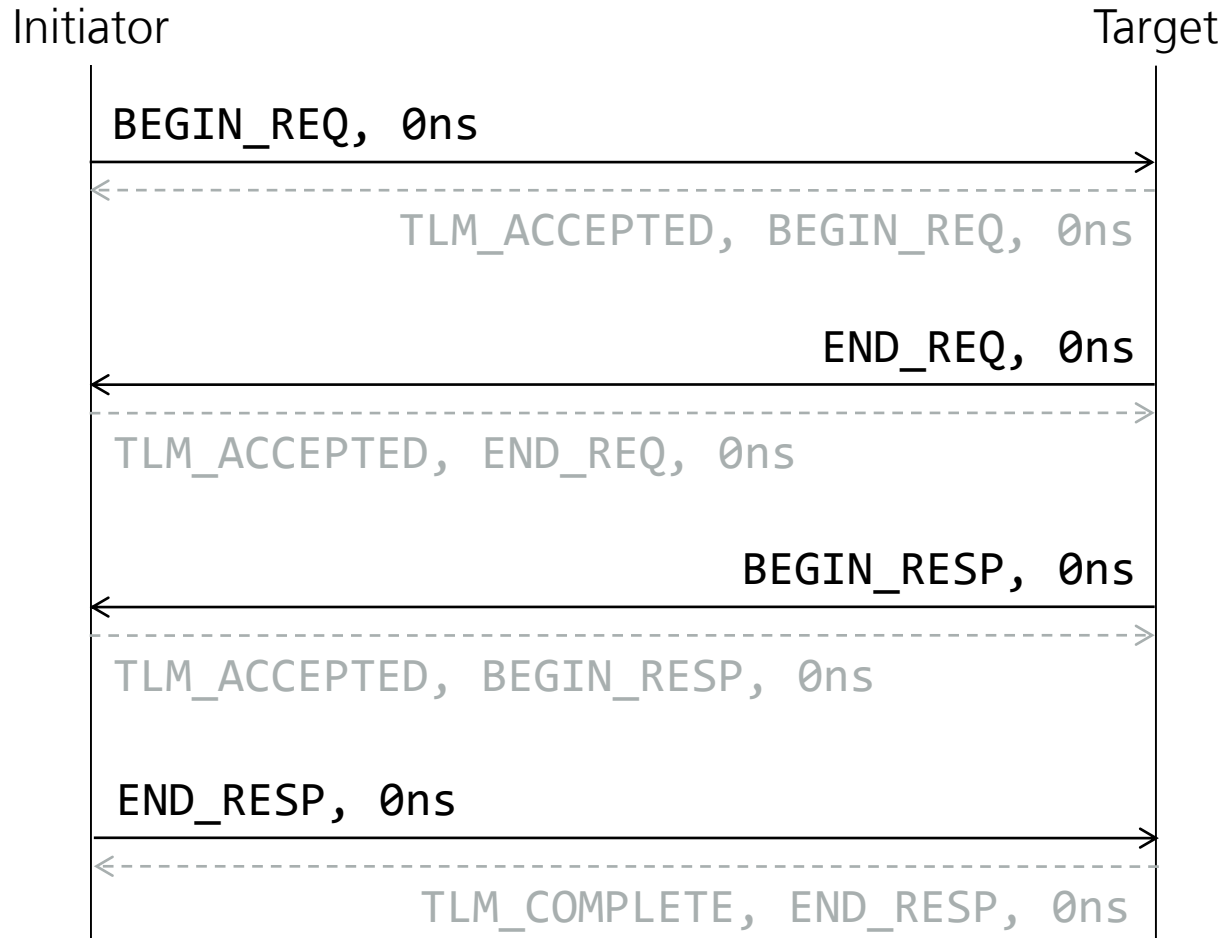
return

Phase transition



29

Base Protocol Rules [1]: 4 Phase Handshake



4 Phase Handshake: Initiator

```
class Initiator: public sc_module, public tlm::tlm_bw_transport_if<> {
public:
    tlm::tlm_initiator_socket<> socket;
    MemoryManager mm;
    int data[16];
    tlm::tlm_generic_payload* requestInProgress;
    sc_event endRequest;
    tlm_utils::peq_with_cb_and_phase<Initiator> peq;

    SC_CTOR(Initiator): socket("socket"),
        requestInProgress(0),
        peq(this, &Initiator::peqCallback)

    {
        socket.bind(*this);
        SC_THREAD(process);
    }

protected:
    void process() {
        tlm::tlm_generic_payload* trans;
        tlm::tlm_phase phase;
        sc_time delay;

        for (int i = 0; i < 100; i++) {
            trans = mm.allocate();
            trans->acquire();

            trans->set_command(...);
            ...
            trans->set_response_status(tlm::TLM_INCOMPLETE_RESPONSE);

            if (requestInProgress) {
                wait(endRequest);
            }
            requestInProgress = trans;
            phase = tlm::BEGIN_REQ;
            delay = ...;

            tlm::tlm_sync_enum status;
            status = socket->nb_transport_fw( *trans, phase, delay );
            ...
            wait(randomDelay());
        }
    }
}
```

Generate a sequence of random transactions

Grab a new transaction from the memory manager pool

BEGIN_REQ/END_REQ exclusion rule

```
virtual tlm::tlm_sync_enum nb_transport_bw(tlm::tlm_generic_payload& trans,
                                           tlm::tlm_phase& phase,
                                           sc_time& delay)
{
    peq.notify(trans, phase, delay);
    return tlm::TLM_ACCEPTED;
}

void peqCallback(tlm::tlm_generic_payload& trans,
                 const tlm::tlm_phase& phase)
{
    if (phase == tlm::END_REQ || ...)
    {
        requestInProgress = 0;
        endRequest.notify();
    }
    else if (phase == tlm::BEGIN_RESP)
    {
        checkTransaction(trans);

        tlm::tlm_phase fw_phase = tlm::END_RESP;
        sc_time delay = sc_time(...);

        socket->nb_transport_fw( trans, fw_phase, delay );

        trans.release();
    }
    else if (phase == tlm::BEGIN_REQ || phase == tlm::END_RESP)
    {
        SC_REPORT_FATAL(name(), "Illegal transaction phase received");
    }
}
};
```

Queue the transaction into the peq until the annotated time has elapsed

Payload event queue callback

Wake-up suspended main process

Do something with transaction

Allow MM to free the transaction object

4 Phase Handshake: Target

HW

```
class Target: public sc_module, public tlm::tlm_fw_transport_if<> {
public:
    tlm::tlm_target_socket<> socket;
    tlm::tlm_generic_payload* transactionInProgress;
    sc_event targetDone;
    bool responseInProgress;
    tlm::tlm_generic_payload* nextResponsePending;
    tlm::tlm_generic_payload* endRequestPending;
    tlm_utils::peq_with_cb_and_phase<Target> peq;

    SC_CTOR(Target) : socket("socket"),
        transactionInProgress(0),
        responseInProgress(false),
        nextResponsePending(0),
        endRequestPending(0),
        peq(this, &Target::peqCallback)
    {
        socket.bind(*this);

        SC_METHOD(executeTransactionProcess);
        sensitive << targetDone; dont_initialize();
    }
    ...
    tlm::tlm_sync_enum nb_transport_fw(tlm::tlm_generic_payload& trans,
        tlm::tlm_phase& phase,
        sc_time& delay)
    {
        peq.notify( trans, phase, delay);
        return tlm::TLM_ACCEPTED;
    }

    void peqCallback(tlm::tlm_generic_payload& trans,
        const tlm::tlm_phase& phase)
    {
        sc_time delay;

        if(phase == tlm::BEGIN_REQ) {
            trans.acquire();

            if (!transactionInProgress) {
                sendEndRequest(trans);
            }
            else {
                endRequestPending = &trans;
            }
        }
    }
}
```

Increment the transaction reference count

Put back-pressure on initiator by deferring END_REQ

```
else if (phase == tlm::END_RESP) {
    ...
    transactionInProgress = 0;
    responseInProgress = false;

    if (nextResponsePending) {
        sendResponse(*nextResponsePending);
        nextResponsePending = 0;
    }

    if (endRequestPending) {
        sendEndRequest(*endRequestPending);
        endRequestPending = 0;
    }
}
else // tlm::END_REQ or tlm::BEGIN_RESP
{
    SC_REPORT_FATAL(name(), "Illegal transaction phase received");
}

void sendEndRequest(tlm::tlm_generic_payload& trans)
{
    tlm::tlm_phase bw_phase;
    sc_time delay;

    bw_phase = tlm::END_REQ;
    delay = ...; // Accept delay

    tlm::tlm_sync_enum status;
    status = socket->nb_transport_bw( trans, bw_phase, delay );

    delay = delay + randomDelay();
    targetDone.notify( delay );

    assert(transactionInProgress == 0);
    transactionInProgress = &trans;
    ...
}
```

Flag must only be cleared when END_RESP is sent

Target itself is now clear to issue the next BEGIN_RESP

unlock the initiator by issuing END_REQ

Queue internal event to mark beginning of response

4 Phase Handshake: Target

HW

```
void executeTransactionProcess()
{
```

```
    executeTransaction(*transactionInProgress);
```

```
    if (responseInProgress) ~
    {
```

```
        ...
```

```
        nextResponsePending = transactionInProgress; ~
```

```
    }
```

```
    else
```

```
    {
```

```
        sendResponse(*transactionInProgress);
```

```
    }
```

```
}
```

```
void sendResponse(tlm::tlm_generic_payload& trans)
```

```
{
```

```
    tlm::tlm_sync_enum status;
```

```
    tlm::tlm_phase bw_phase;
```

```
    sc_time delay;
```

```
    responseInProgress = true;
```

```
    bw_phase = tlm::BEGIN_RESP;
```

```
    delay = SC_ZERO_TIME;
```

```
    status = socket->nb_transport_bw( trans, bw_phase, delay );
```

```
    ...
```

```
    trans.release();
```

```
}
```

```
...
```

```
};
```

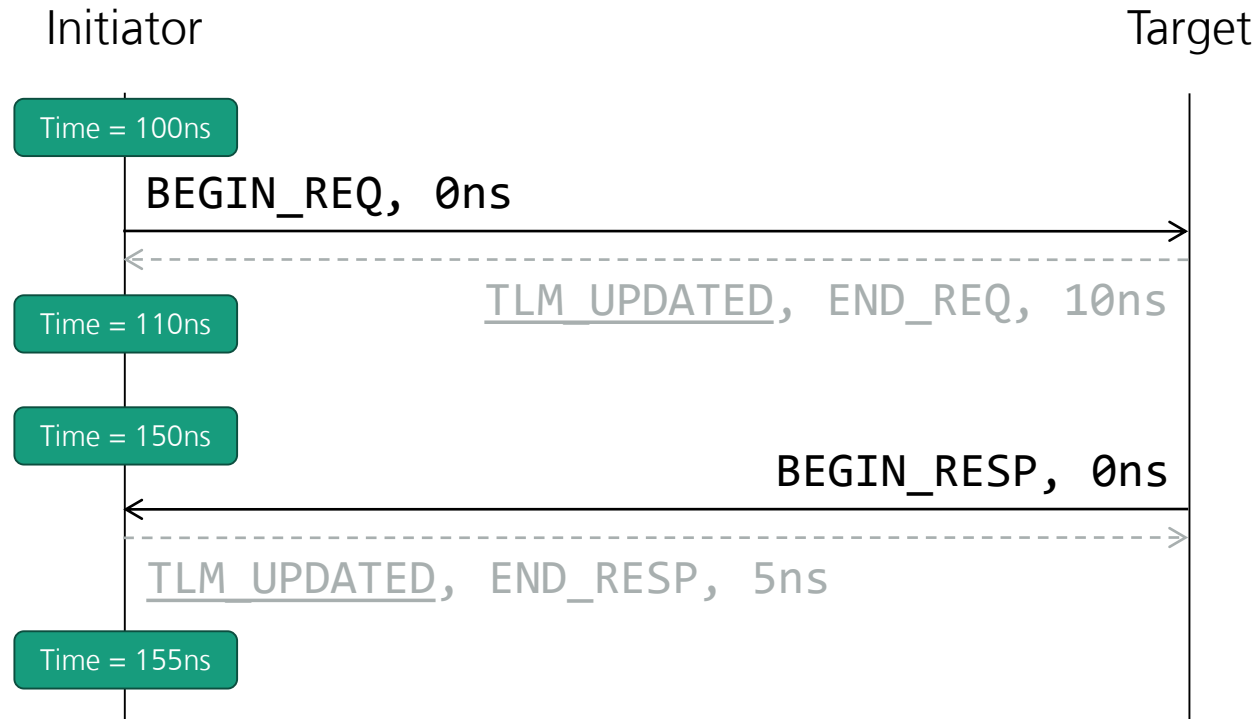
Method process that runs on
targetDone event

Target must honor
BEGIN_RESP/END_RE
SP exclusion rule

Function for sending
the response

Tell the memory
manager to free
transaction

Base Protocol Rules [2]: Using the Return Paths



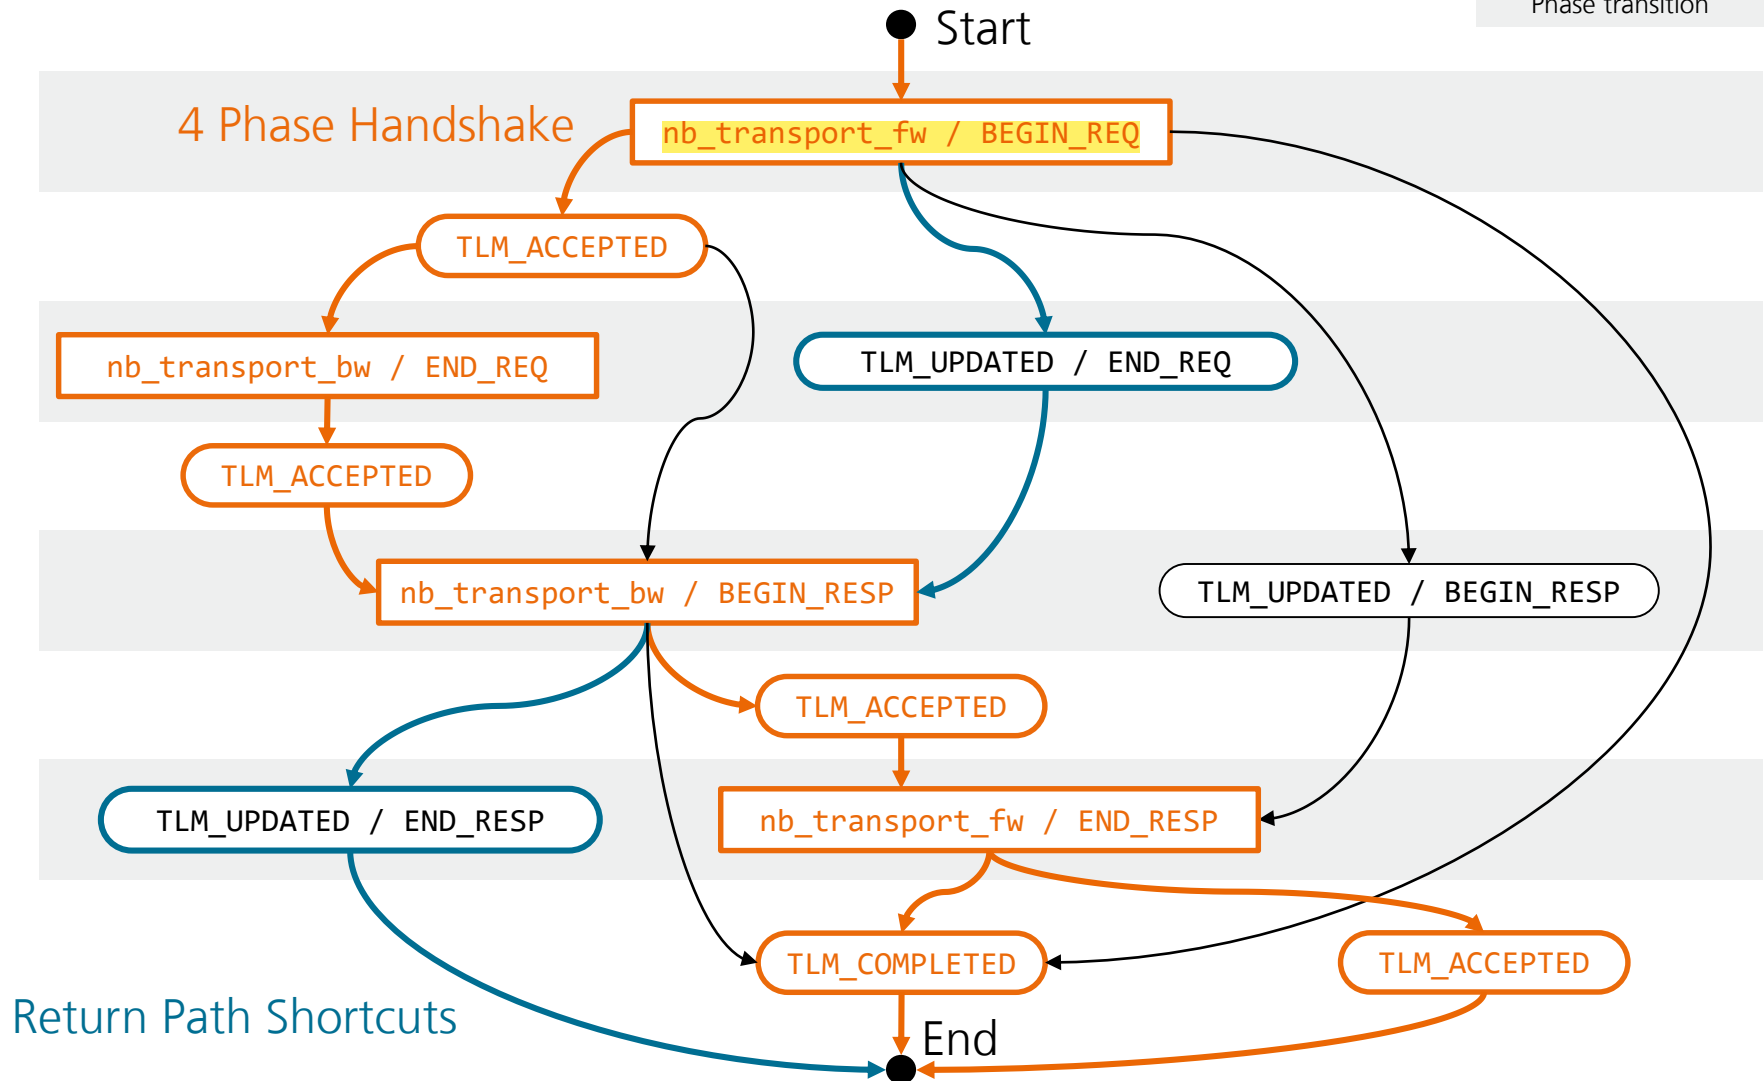
- The FW return path can be used as an alternative to the BW path
- The BW return path can be used as an alternative to the FW path
 - State must be set to **TLM_UPDATED**
- Timing must be increased

Full Overview of AT Protocol Possibilities

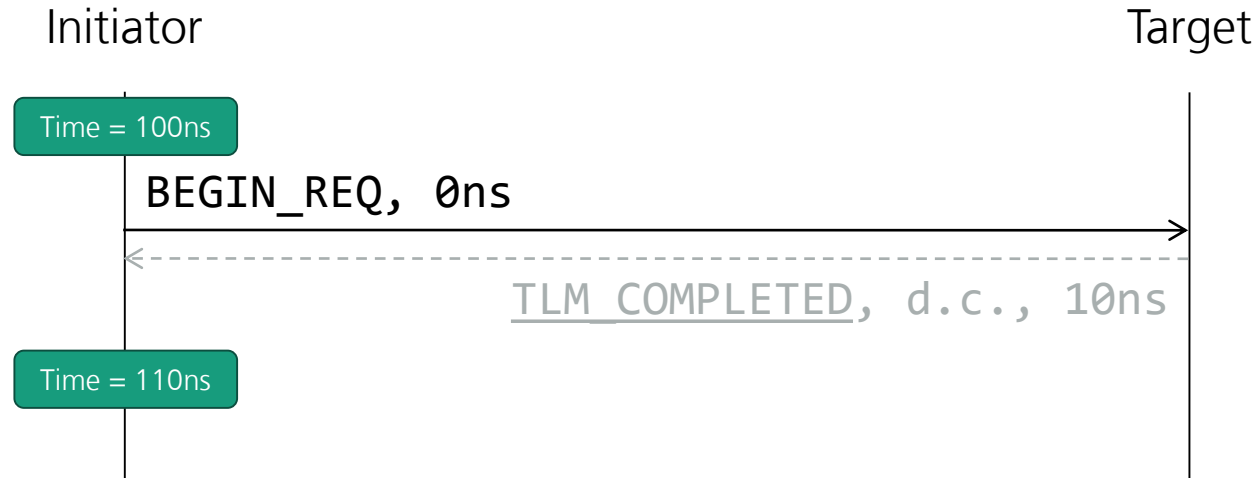
call

return

Phase transition

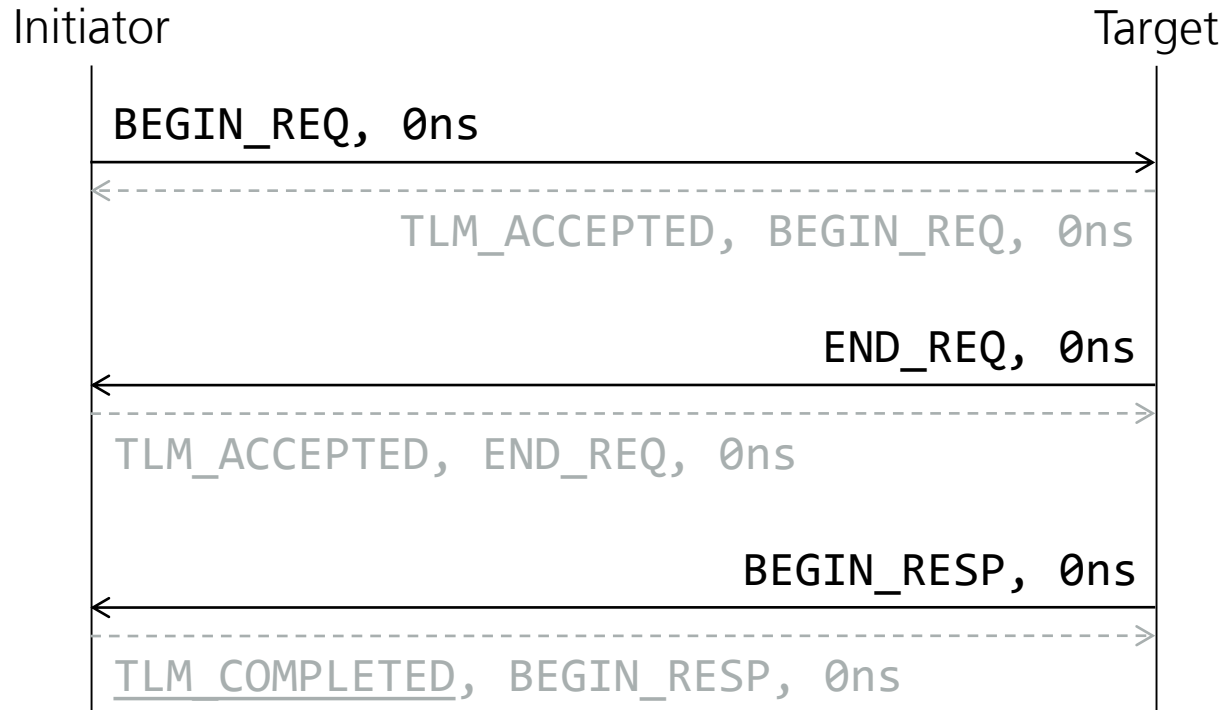


Base Protocol Rules [4]: Early Completion



- The initiator sends **BEGIN_REQ** and the target immediately returns **TLM_COMPLETED** (useful for modelling simple I/O, where no ACK is required)
- The targets pre-empt any further communication by signaling a so called *Early completion*
- Initiator should immediacy check the response status of the generic payload object
- With **TLM_COMPLETED** the caller should always ignore the phase argument

Base Protocol Rules [4]: Early Completion



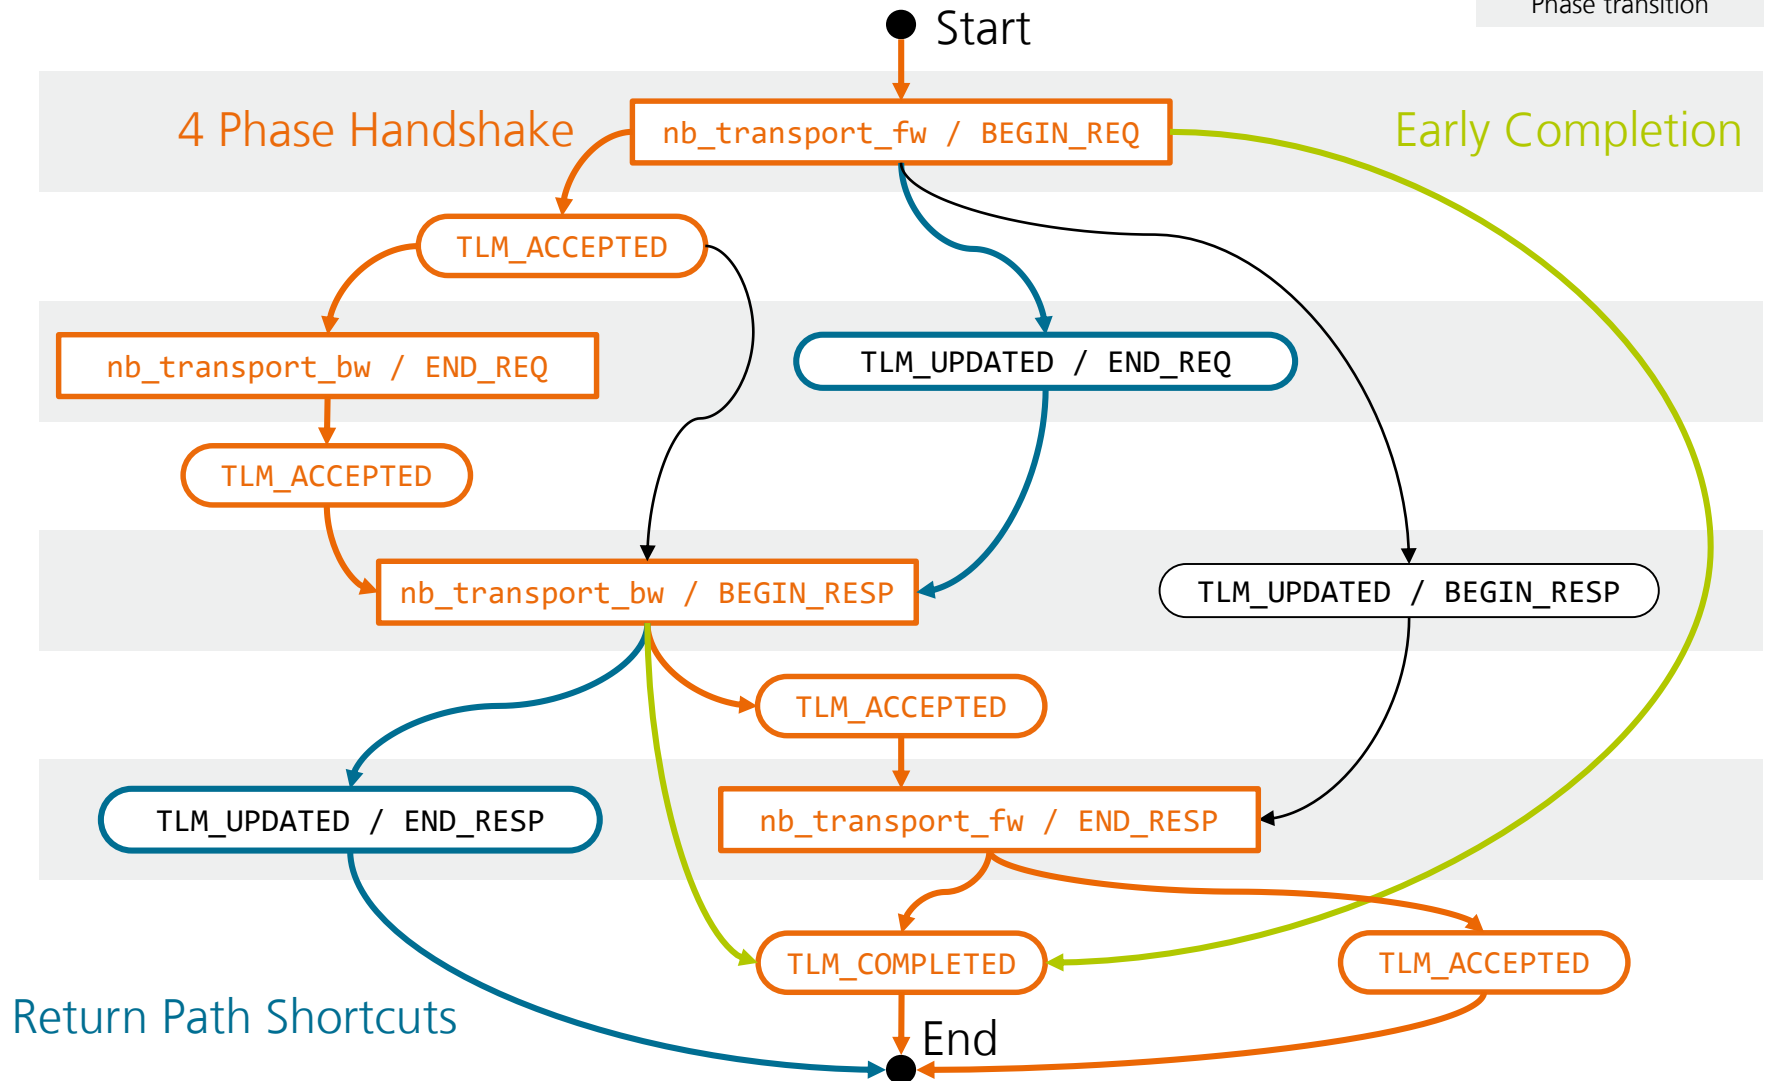
- The initiator can early complete the transaction by returning `TLM_COMPLETED`
- The initiator pre-empts therefore any further communication

Full Overview of AT Protocol Possibilities

call

return

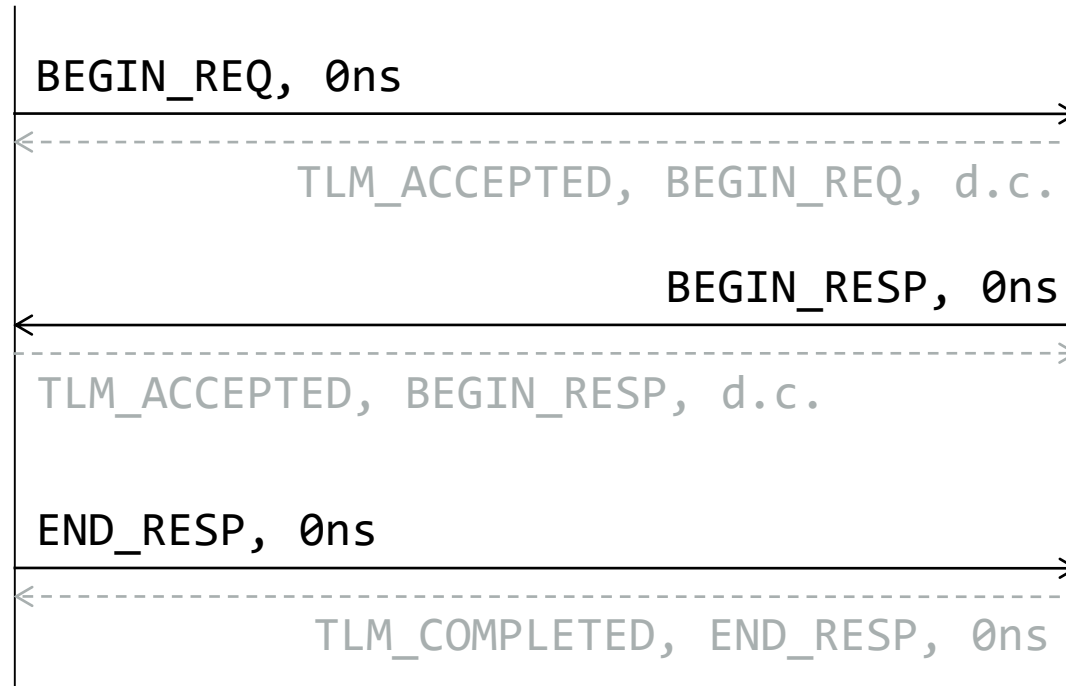
Phase transition



Base Protocol Rules [3]: Skip END_REQ

Initiator

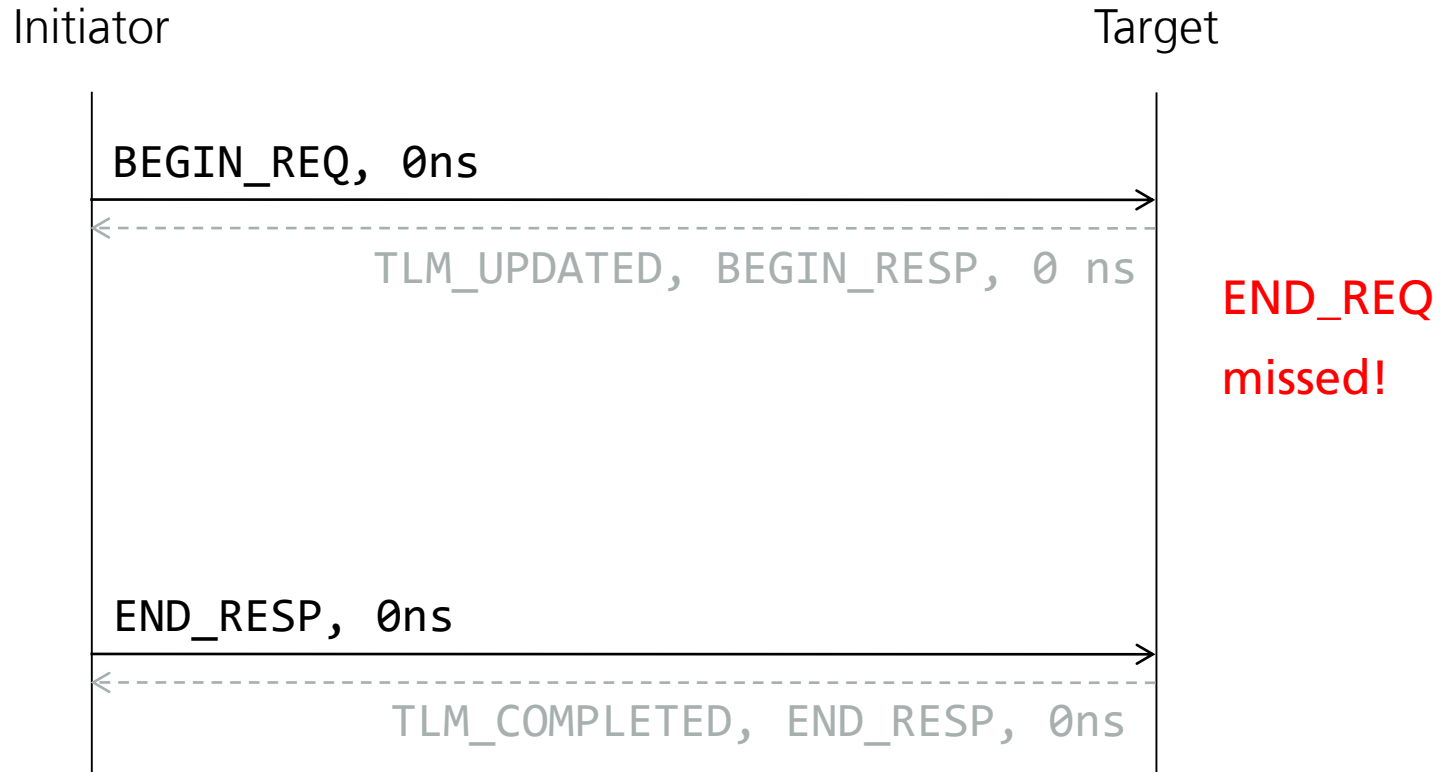
Target



**END_REQ
missed!**

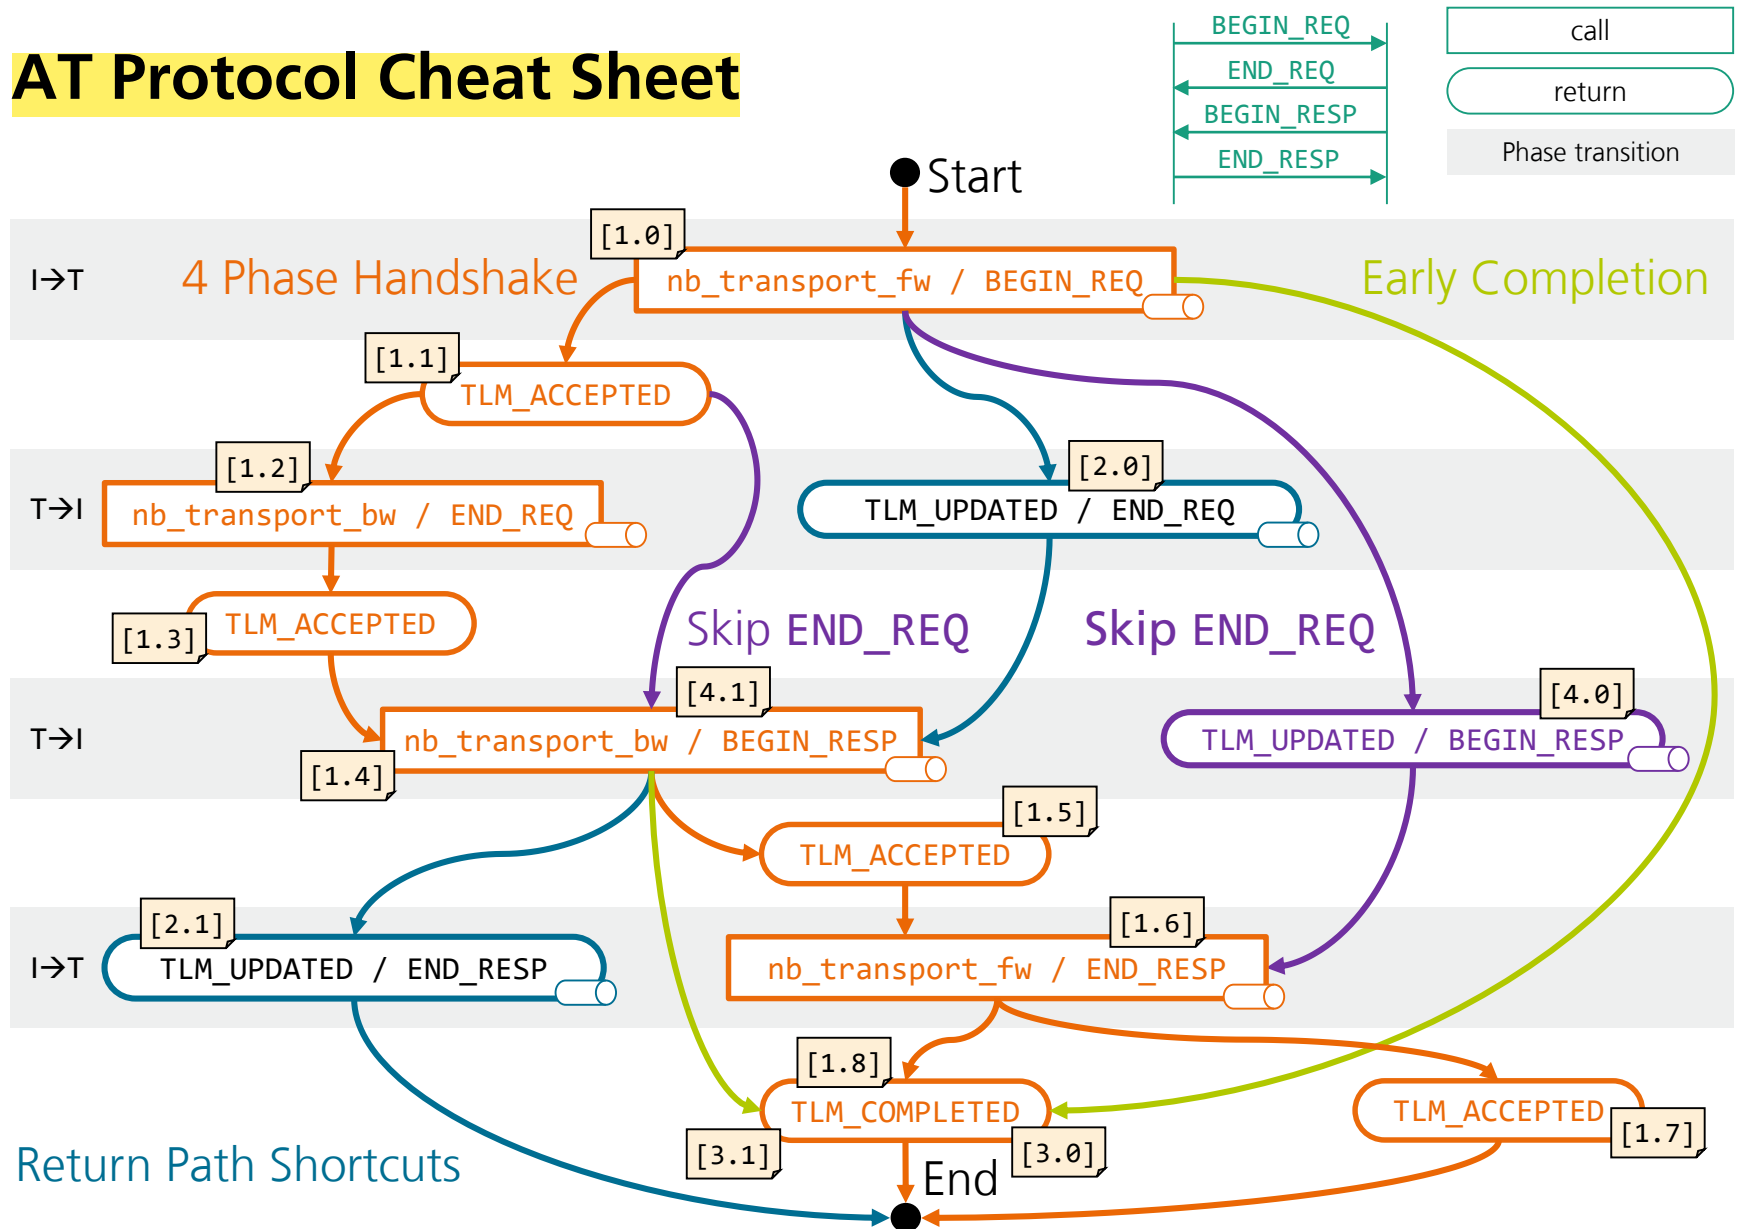
- END_REQ is pre-empted by the target sending BEGIN_RESP in the next BW call

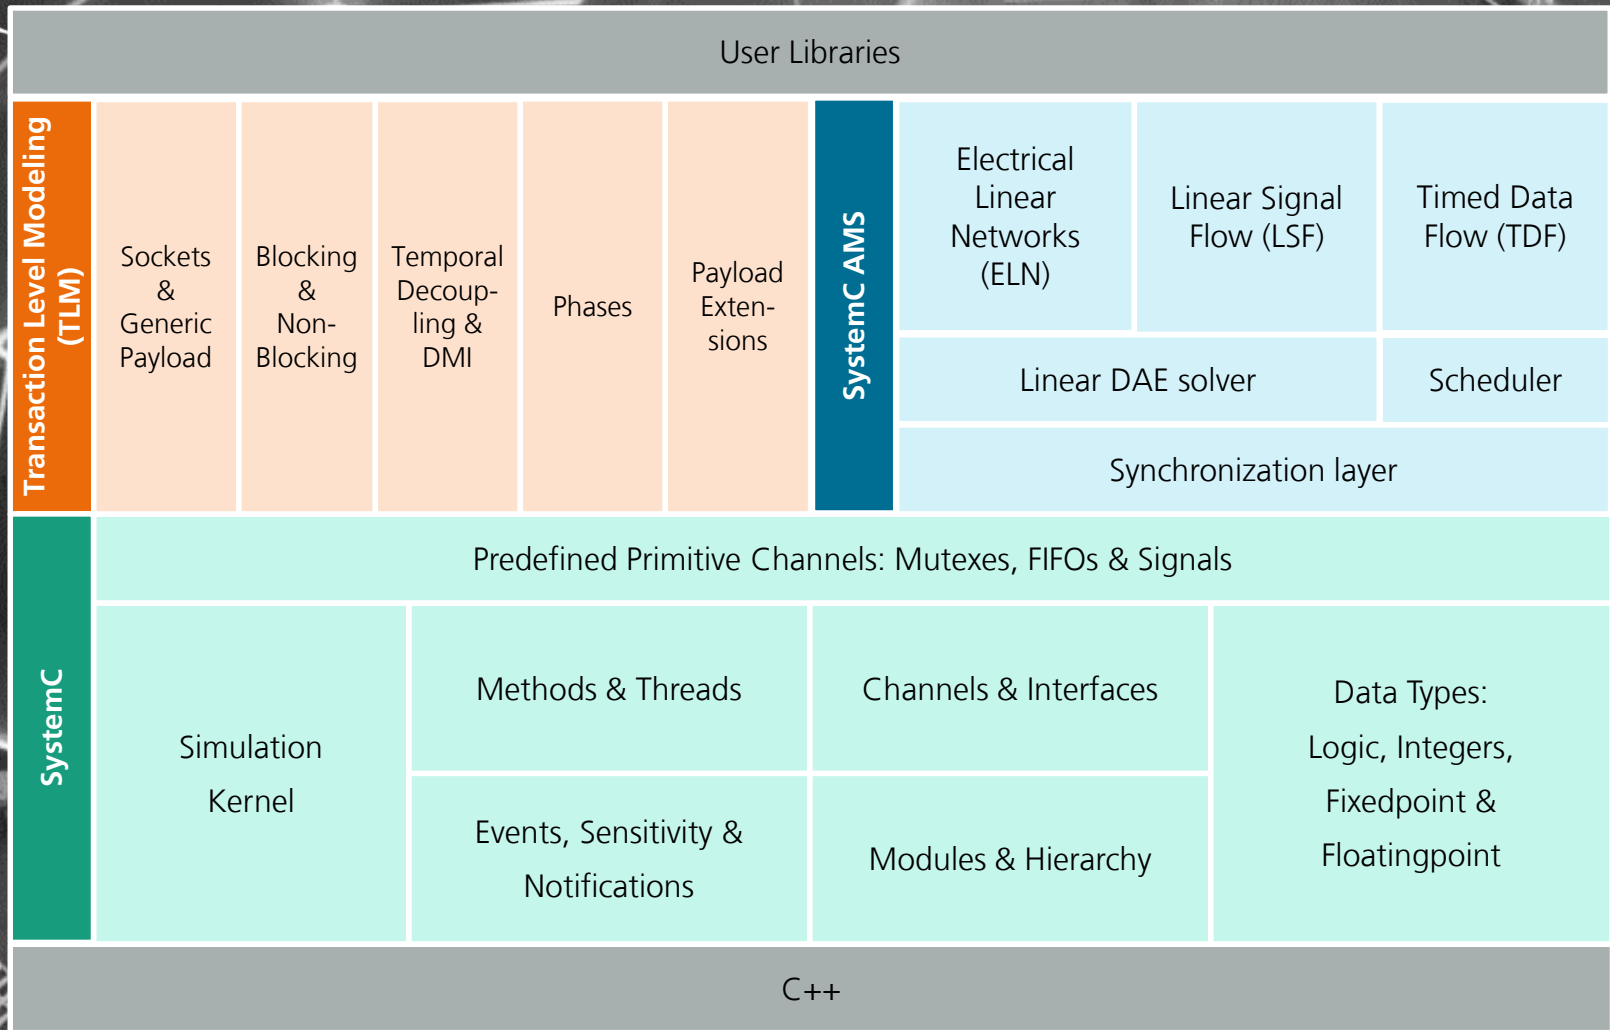
Base Protocol Rules (7): Skip END_REQ (Shortcut)



- END_REQ is pre-empted by the target sending directly BEGIN_RESP via TLM_UPDATED over the return path

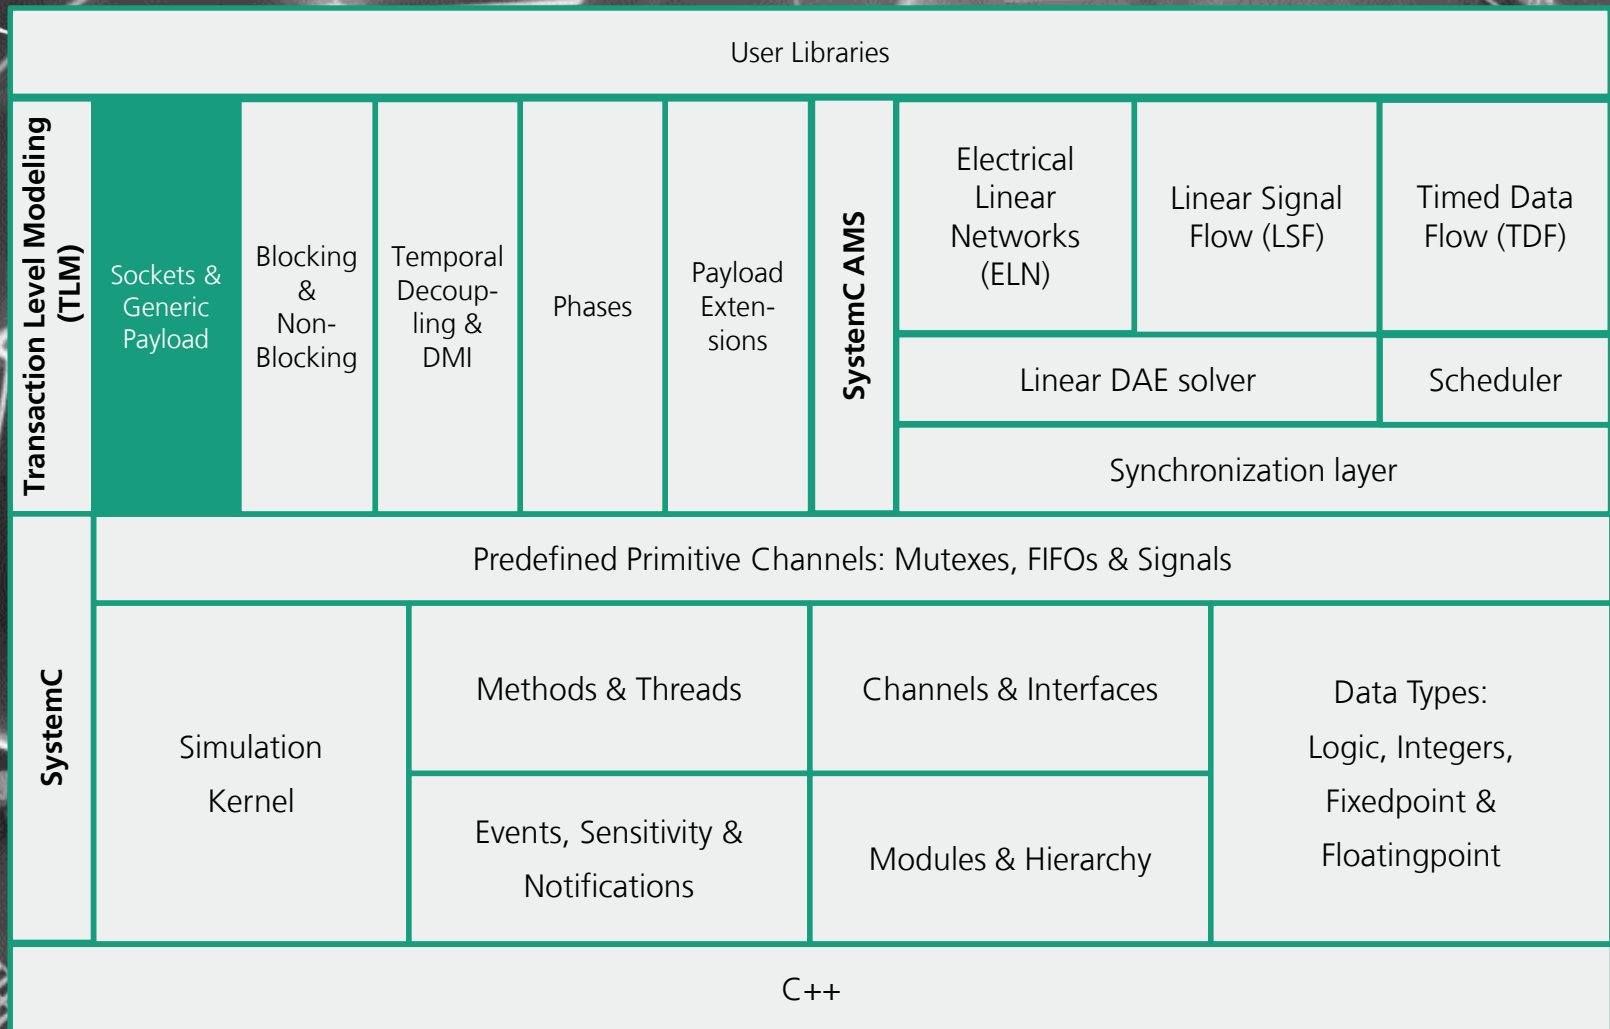
AT Protocol Cheat Sheet







TLM Advanced



Simple Sockets

- “Simple” because they are easy to use – less to code ...
- Do not bind sockets to objects, instead register methods with each socket
- Dummy method implementations are provided:
 - `get_direct_mem_ptr`
 - `transport_dbg`
 - `invalidate_direct_mem_ptr`
 - `nb_transport_bw`
- Target need only register either `b_transport` or `nb_transport_fw`
- Automatic conversion between blocking and non-blocking

Simple Sockets: Initiator Example

```
#include "tlm_utils/simple_initiator_socket.h"
...

SC_MODULE(Initiator) {
    public:
        tlm_utils::simple_initiator_socket<Initiator> iSocket;
        ...
    public:
        SC_CTOR(Initiator): iSocket("iSocket"), requestInProgress(0), peq(this, &Initiator::peqCallback)
        {
            iSocket.register_nb_transport_bw(this, &Initiator::nb_transport_bw);

            SC_THREAD(process);
            ...
        }

        void process() {
            ...
        }

        tlm::tlm_sync_enum nb_transport_bw(tlm::tlm_generic_payload& ..., tlm::tlm_phase& ..., sc_time& ...) {
            ...
        }
        ...
};
```

Instantiate simple socket

Register function

Implementation of Function

Try code on github:
https://github.com/TUK-SCVP/SCVP.artifacts/blob/master/tlm_simple_sockets/

Simple Sockets: Target Example

```
#include "tlm_utils/simple_target_socket.h"
...
SC_MODULE(Target)
{
    tlm_utils::simple_target_socket<Target> tSocket;
    ...
    SC_CTOR(Target) : tSocket("tSocket"), ...
    {
        tSocket.register_b_transport(this, &Target::b_transport);
        tSocket.register_nb_transport_fw(this, &Target::nb_transport_fw);
    }

    void b_transport(tlm::tlm_generic_payload& trans, sc_time& delay) {
        ...
    }

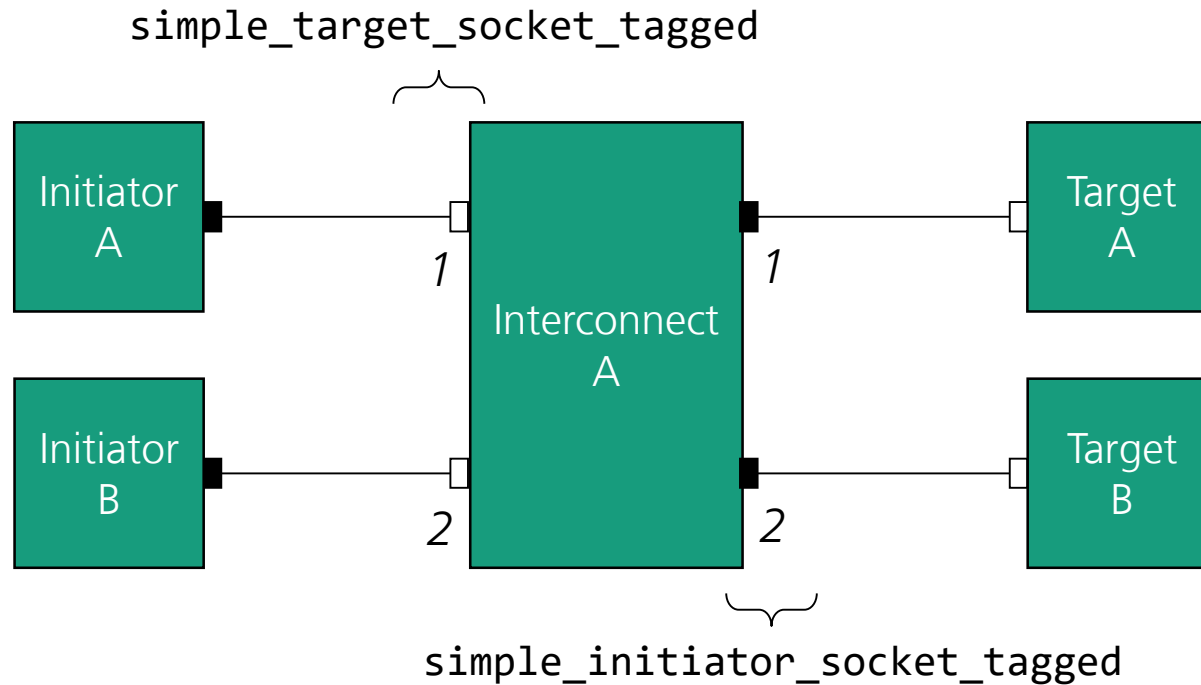
    tlm_sync_enum nb_transport_fw(tlm::tlm_generic_payload& ..., tlm::tlm_phase& ..., sc_time& ...) {
        ...
    }
    ...
};
```

Instantiate simple
target socket

Register functions

Implementation of
Functions

Tagged Simple Sockets: Interconnect Example



- Register the same callback method with several sockets
- Distinguish origin of incoming transactions using socket id
- Interconnect is usually template class for number of target and initiator sockets

Tagged Simple Sockets: Interconnect Example

```
template<unsigned int I, unsigned int T>
SC_MODULE(BUS)
{
    public:
        tlm_utils::simple_target_socket_tagged<BUS> tSocket[T];
        tlm_utils::simple_initiator_socket_tagged<BUS> iSocket[I];

        SC_CTOR(BUS)
        {
            for(unsigned int i = 0; i < T; i++)
            {
                tSocket[i].register_b_transport(this,
                    &BUS::b_transport, i);
                tSocket[i].register_nb_transport_fw(this,
                    &BUS::nb_transport_fw, i);
            }

            for(unsigned int i = 0; i < I; i++)
            {
                iSocket[i].register_nb_transport_bw(this,
                    &BUS::nb_transport_bw, i);
            }
        }

        private:
            std::map<tlm::tlm_generic_payload*, int> bwRoutingTable;
            std::map<tlm::tlm_generic_payload*, int> fwRoutingTable;

            virtual void b_transport( int id,
                                     tlm::tlm_generic_payload& trans,
                                     sc_time& delay )
            {
                {
                    sc_assert(id < T);
                    int outPort = routeFW(id, trans, false);
                    iSocket[outPort]->b_transport(trans, delay);
                }
            }
        }
    };

```

```
virtual tlm::tlm_sync_enum nb_transport_fw( int id,
                                             tlm::tlm_generic_payload& trans,
                                             tlm::tlm_phase& phase,
                                             sc_time& delay )
{
    sc_assert(id < T);
    int outPort = 0;

    if(phase == tlm::BEGIN_REQ) {
        trans.acquire();
        outPort = routeFW(id, trans, true);
    }
    else if(phase == tlm::END_RESP) {
        outPort = fwRoutingTable[&trans];
        trans.release();
    }
    else {
        SC_REPORT_FATAL(name(), "ERROR!");
    }

    return iSocket[outPort]->nb_transport_fw(
        trans, phase, delay);
}

virtual tlm::tlm_sync_enum nb_transport_bw( int id,
                                             tlm::tlm_generic_payload& trans,
                                             tlm::tlm_phase& phase,
                                             sc_time& delay )
{
    int inPort = bwRoutingTable[&trans];

    return tSocket[inPort]->nb_transport_bw(
        trans, phase, delay);
}
};

```

Simple Sockets: Target Example

```
int routeFW(int inPort, tlm::tlm_generic_payload &trans, bool store)
{
    int outPort = 0;

    if(trans.get_address() < 512)
    {
        outPort = 0;
    }
    else if(trans.get_address() >= 512 && trans.get_address() < 1024)
    {
        trans.set_address(trans.get_address() - 512);
        outPort = 1;
    }
    else {
        trans.set_response_status( tlm::TLM_ADDRESS_ERROR_RESPONSE );
    }

    if(store) {
        bwRoutingTable[&trans] = inPort;
        fwRoutingTable[&trans] = outPort;
    }

    return outPort;
}
```

Implementation of
the memory map

Correct address,
i.e. subtract offset
such that target
gets always
addresses starting
at 0

Store from where
transaction comes
and where it goes

Simple Sockets: Target Example

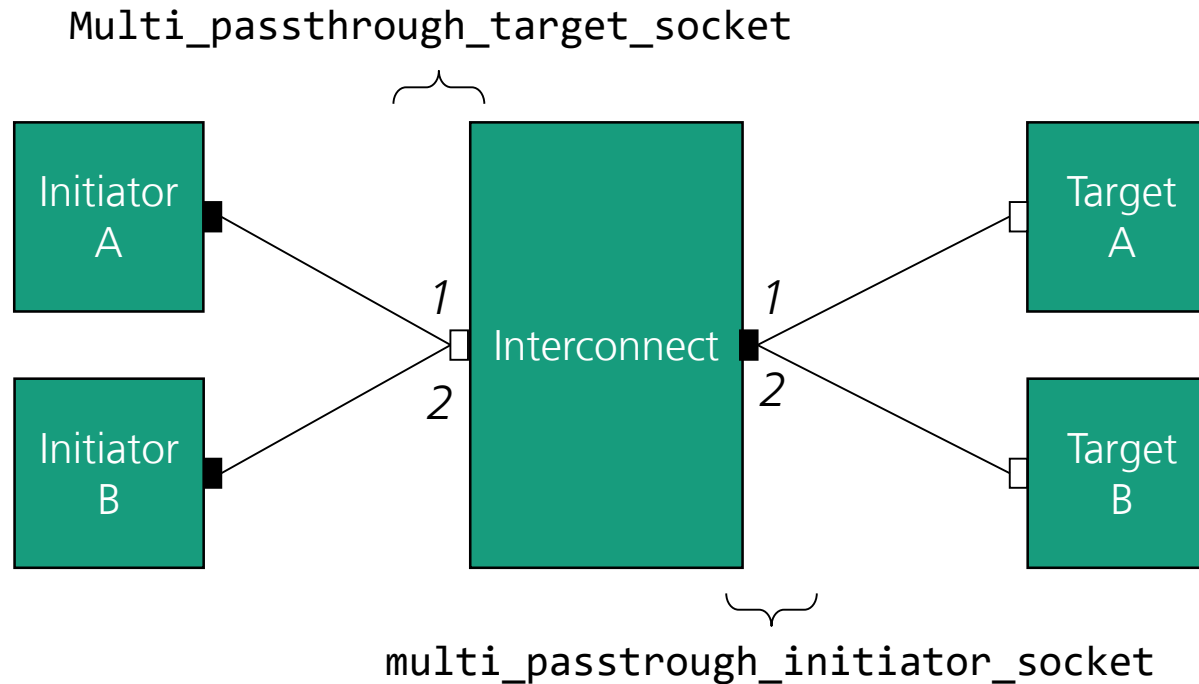
```
int sc_main (...)  
{  
  
    Initiator * cpu1    = new Initiator("C1");  
    Initiator * cpu2    = new Initiator("C2");  
  
    Target * memory1    = new Target("M1");  
    Target * memory2    = new Target("M2");  
  
    Interconnect<2,2> * bus = new BUS<2,2>("B1");  
  
    cpu1->iSocket.bind(bus->tSocket[0]);  
    cpu2->iSocket.bind(bus->tSocket[1]);  
  
    bus->iSocket[0].bind(memory1->tSocket);  
    bus->iSocket[1].bind(memory2->tSocket);  
  
    sc_start();  
  
    return 0;  
}
```

Number of components must be known at compile time!

Try code on github:

https://github.com/TUK-SCVP/SCVP.artifacts/blob/master/tlm_simple_sockets/

Multipassthrough Sockets



- Similar to Tagged sockets: also an id is used for identification
- The id is determined by the order of the binding to the multipassthrough socket
- Dynamic binding, number of comp. does not have to be known at compile time

Multipassthrough Sockets: Interconnect Example

```
SC_MODULE(BUS)
{
    public:
        tlm_utils::multi_passthrough_target_socket<BUS> tSocket;
        tlm_utils::multi_passthrough_initiator_socket<BUS> iSocket;

        SC_CTOR(Interconnect) : tSocket("tSocket"), iSocket("iSocket")
        {
            tSocket.register_b_transport(this, &BUS::b_transport);
            tSocket.register_nb_transport_fw(this, &BUS::nb_transport_fw);
            iSocket.register_nb_transport_bw(this, &BUS::nb_transport_bw);
        }

        private:
            std::map<tlm::tlm_generic_payload*, int> bwRoutingTable;
            std::map<tlm::tlm_generic_payload*, int> fwRoutingTable;

            void b_transport( int id, tlm::tlm_generic_payload& trans, sc_time& delay ) {
                ...
            }

            tlm::tlm_sync_enum nb_transport_fw( int id, ... ) {
                ...
            }

            virtual tlm::tlm_sync_enum nb_transport_bw( int id, ... ) {
                ...
            }
};
```

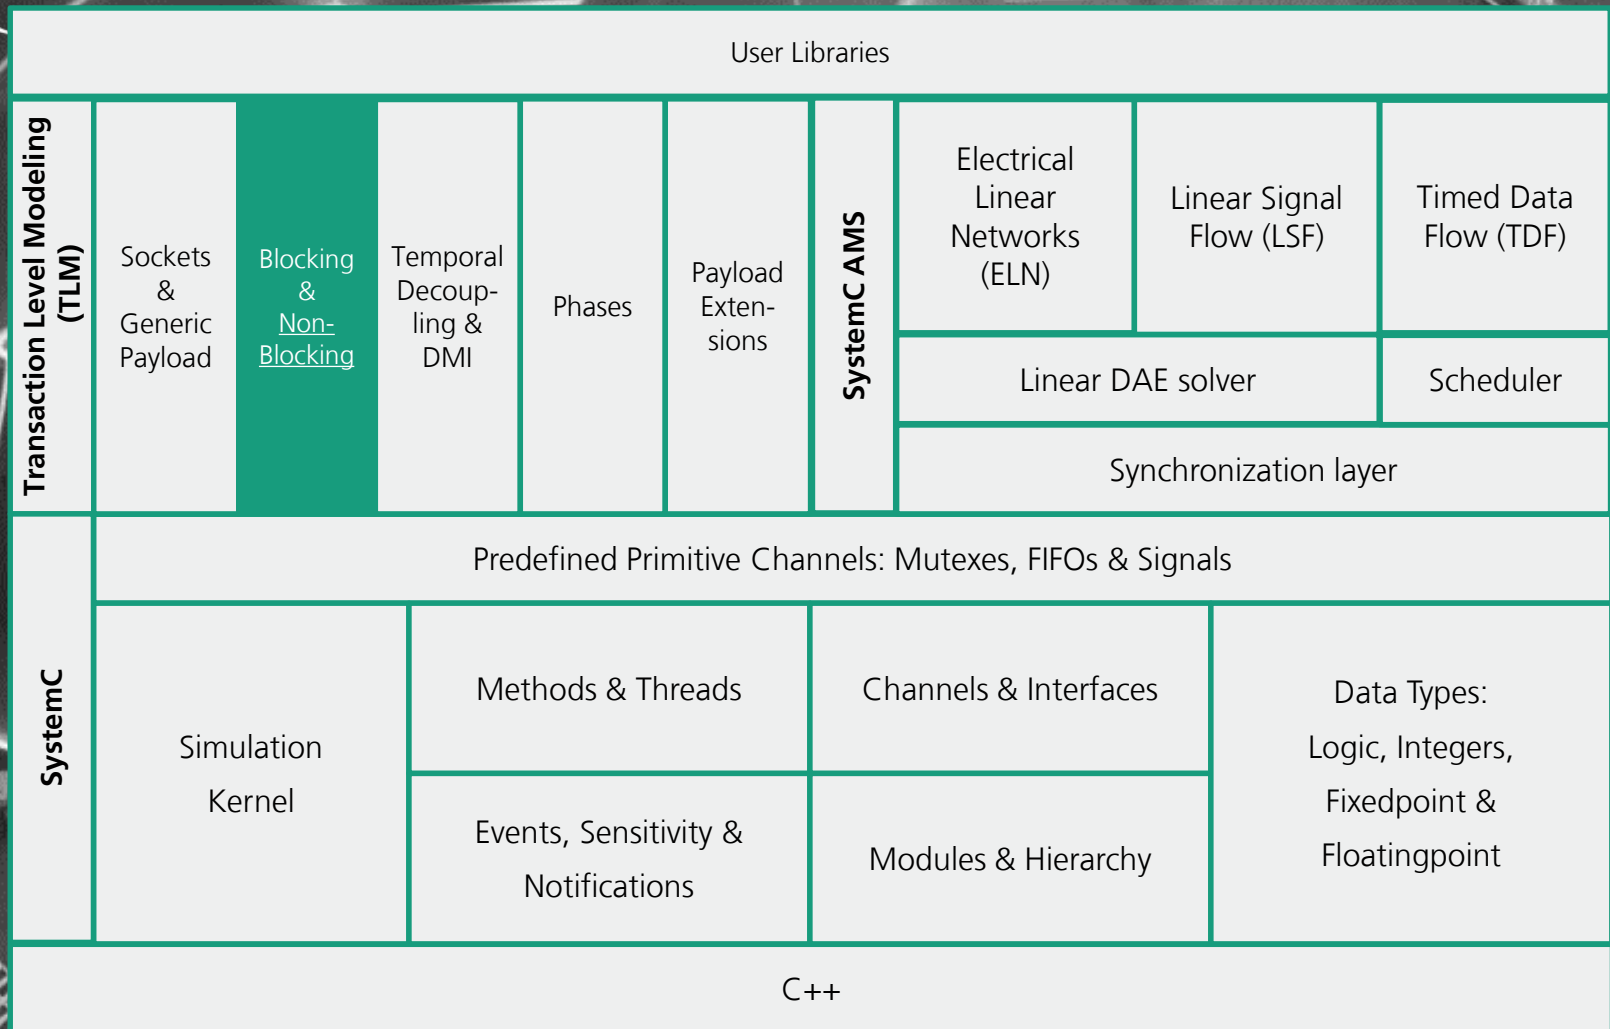
Multipasstrough Sockets: Interconnect Example

```
int sc_main (...)  
{  
    Initiator * cpu1    = new Initiator("C1");  
    Initiator * cpu2    = new Initiator("C2");  
  
    Target * memory1    = new Target("M1");  
    Target * memory2    = new Target("M2");  
  
    Interconnect * bus = new BUS("B1");  
  
    cpu1->iSocket.bind(bus->tSocket);  
    cpu2->iSocket.bind(bus->tSocket);  
  
    bus->iSocket.bind(memory1->tSocket);  
    bus->iSocket.bind(memory2->tSocket);  
  
    sc_start();  
  
    return 0;  
}
```

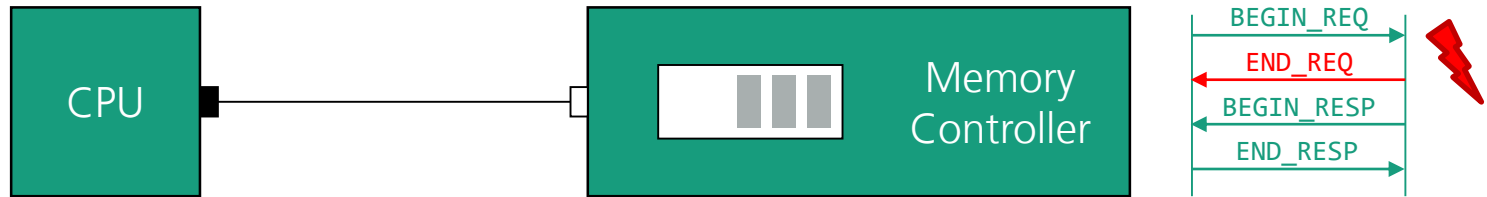
The order of the binding determines the ids for the function calls!

Try code on github:

https://github.com/TUK-SCVP/SCVP.artifacts/blob/master/tlm_multipasstrough_sockets



Modelling of Backpressure



- The exclusion rules enable flow control like back pressure:
 - E.g. if an input buffer of a target is full the target can defer the sending of **END_REQ** for the transaction that filled the buffer, until the buffer has available space again
 - An RTL ready signal can be modeled by deferring **END_REQ**
 - However, since it is non-blocking, the target can do something else, e.g. sending
- Also Response exclusion rule must be honored!

Modelling of Backpressure

HW

```
DECLARE_EXTENDED_PHASE(INTERNAL);
```

```
SC_MODULE(Target) {
```

```
    bool responseInProgress;  
    tlm::tlm_generic_payload* endRequestPending;  
    tlm_utils::peq_with_cb_and_phase<Target> peq;  
    unsigned int numberOfTransactions;  
    unsigned int bufferSize;  
    std::queue<tlm::tlm_generic_payload*> responseQueue;  
    ...
```

```
    void peqCallback(...) {  
        sc_time delay;
```

```
        if(phase == tlm::BEGIN_REQ) {  
            trans.acquire();  
            if (numberOfTransactions < bufferSize) {  
                sendEndRequest(trans);  
            } else {  
                endRequestPending = &trans;  
            }  
        }
```

```
    } else if (phase == tlm::END_RESP) {  
        ...  
        numberOfTransactions--;  
        ...  
        if (responseQueue.size() > 0) {  
            gp * next = responseQueue.front();  
            responseQueue.pop();  
            sendResponse(*next);  
        }  
        if (endRequestPending) {  
            sendEndRequest(*endRequestPending);  
            endRequestPending = 0;  
        }  
    }
```

```
    else if(phase == INTERNAL) {  
        executeTransaction(trans);  
        if (responseInProgress) {  
            responseQueue.push(&trans);  
        } else {  
            sendResponse(trans);  
        }  
    }  
}
```

Internal protocol phase, instead of using a sensitive process!

Check if input buffer is full

Reduce transaction counter

Send next response

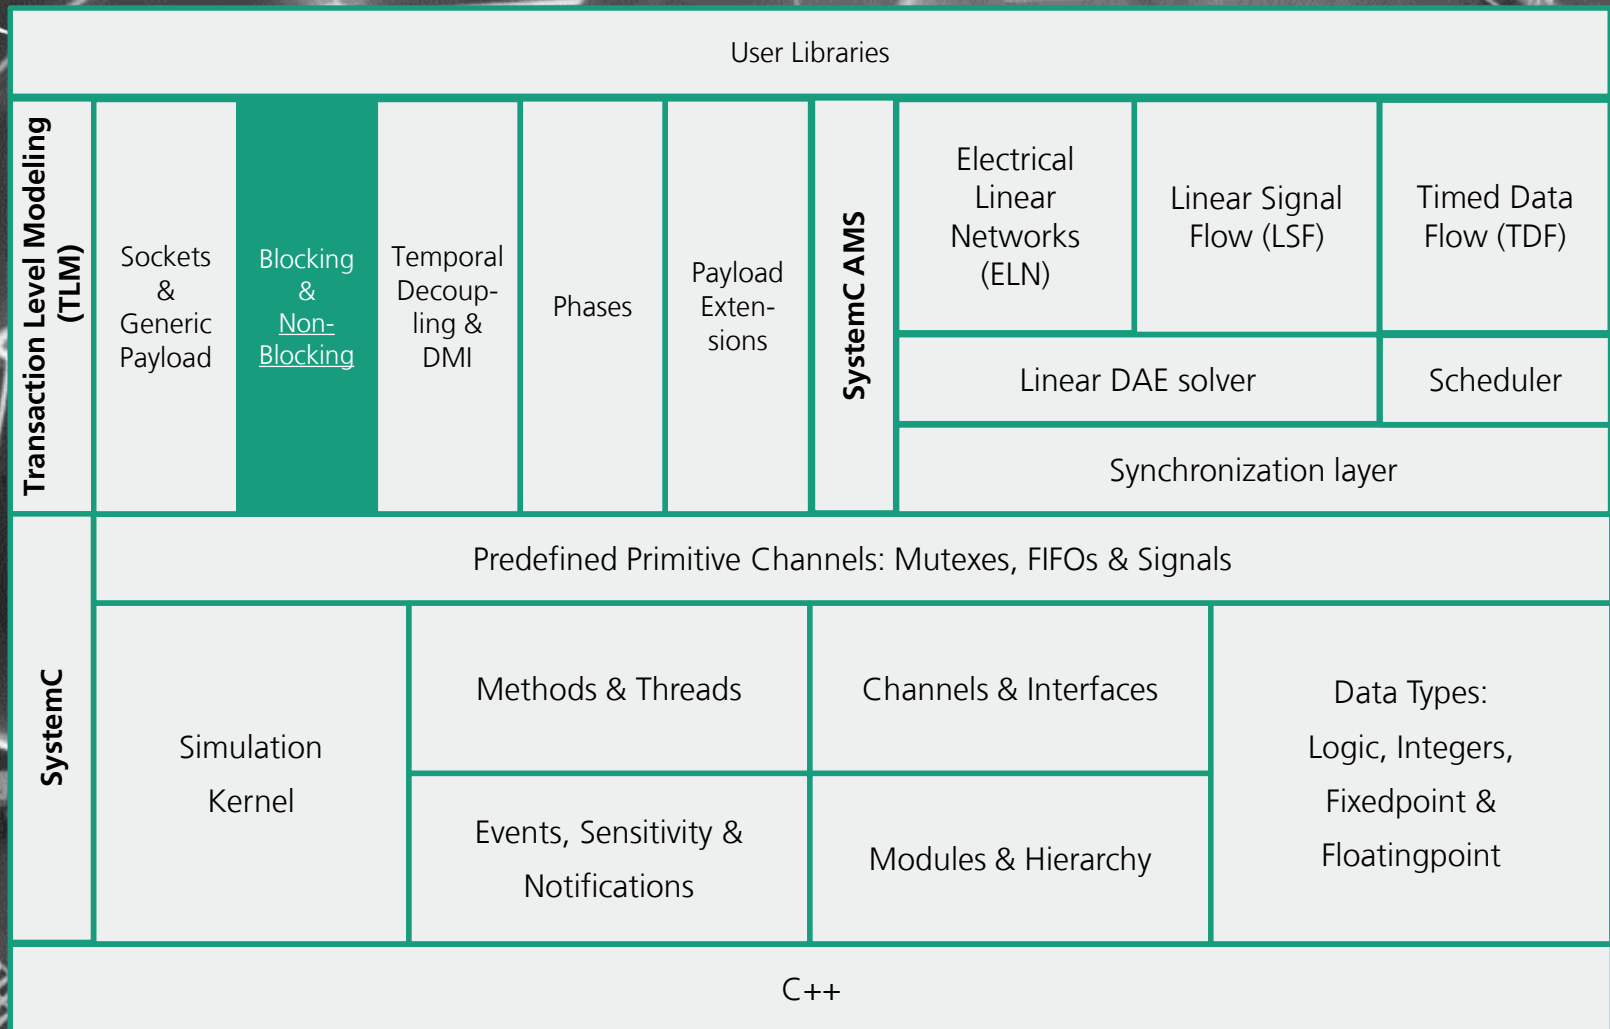
Unblock Initiator

Honor response exclusion rule

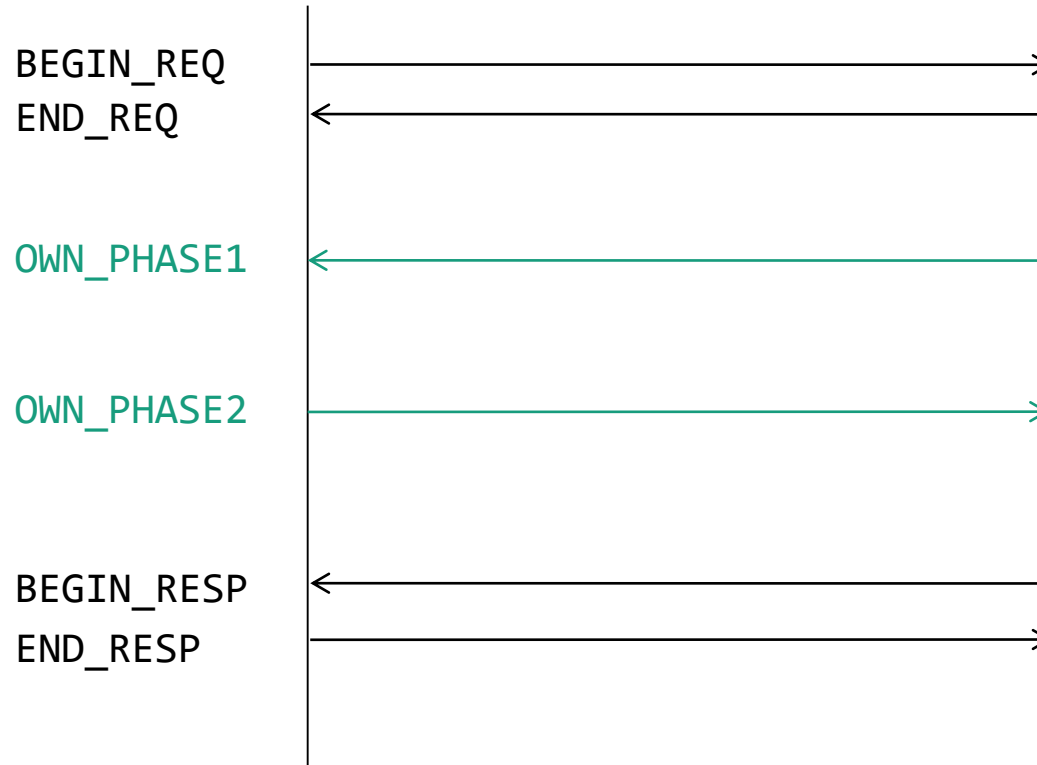
```
void sendEndRequest(tlm::tlm_generic_payload& trans) {  
    tlm::tlm_phase bw_phase;  
    sc_time delay;  
    bw_phase = tlm::END_REQ;  
    delay = randomDelay();  
  
    tlm::tlm_sync_enum status;  
    status = socket->nb_transport_bw(trans, bw_phase, delay);  
    delay = delay + randomDelay();  
    peq.notify(trans, INTERNAL, delay);  
  
    numberOfTransactions++;  
    printBuffer(bufferSize, numberOfTransactions);  
}  
  
void sendResponse(tlm::tlm_generic_payload& trans) {  
    tlm::tlm_sync_enum status;  
    tlm::tlm_phase bw_phase;  
    sc_time delay;  
  
    responseInProgress = true;  
    bw_phase = tlm::BEGIN_RESP;  
    delay = SC_ZERO_TIME;  
    status = socket->nb_transport_bw(trans, bw_phase, delay);  
  
    if (status == tlm::TLM_UPDATED) {  
        peq.notify(trans, bw_phase, delay);  
    }  
    else if (status == tlm::TLM_COMPLETED) {  
        numberOfTransactions--;  
        responseInProgress = false;  
    }  
    trans.release();  
}  
};
```

Try code on github:

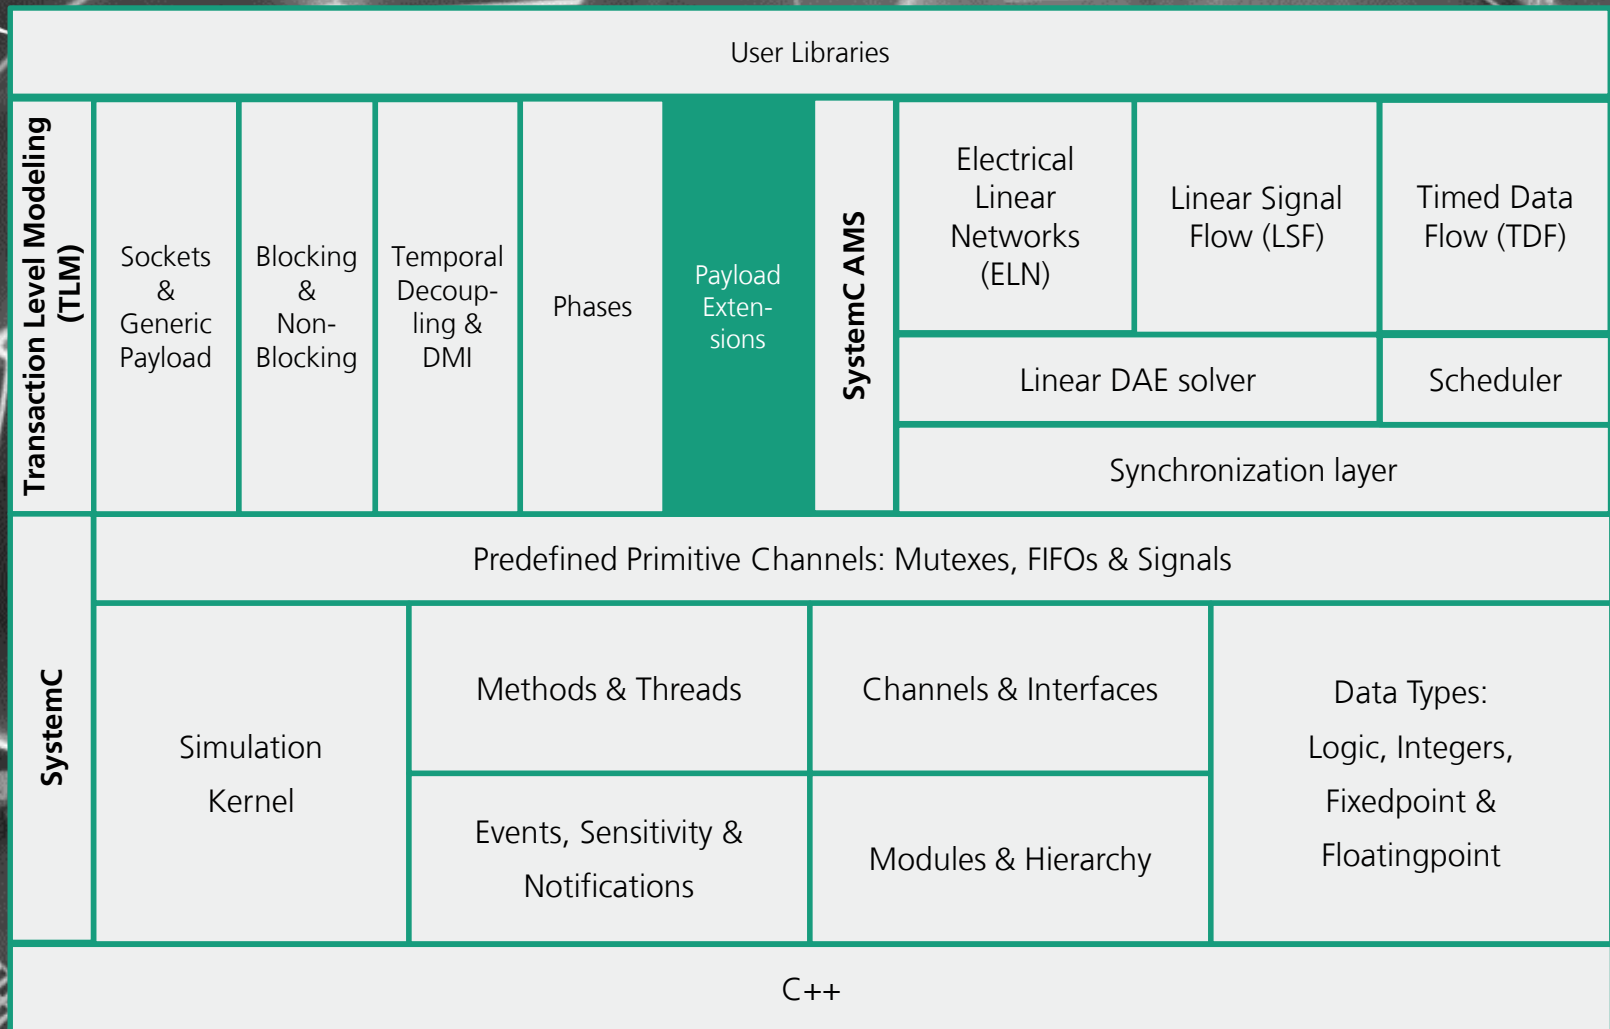
https://github.com/TUK-SCVP/SCVP.artifacts/blob/master/tlm_at_backpressure



Self defined Protocol Phases



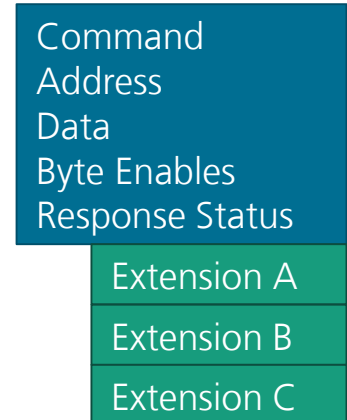
- Base protocol can be enhanced by phase extensions
- Usage of the macro: **DECLARE_EXTENDED_PHASE**(INTERNAL);
- More synchronization points → higher accuracy → slower simulation speed



Payload Extensions

- Payload extensions are a flexible and powerful method to carry additional non-standard attributes (e.g. priority, ...)
- Extensions are attached to the normal generic payload
- Extensions can be:
 - Ignorable (ensures compatibility with base protocol)
 - Mandatory (a new protocol is needed!)
 - Private (only used by a single module e.g. Interconnect for storing routing information instead using maps)
 - Sticky (remain when transaction is returned to a pool)
 - Auto (freed when transaction is returned to a pool)

Generic payload object



Payload Extensions Example

```
class routingExtension : public
tlm::tlm_extension<routingExtension> {
private:
    unsigned int inputPortNumber;
    unsigned int outputPortNumber;

public:
    routingExtension(unsigned int i, unsigned int o) :
        inputPortNumber(i), outputPortNumber(o)
    {}

    tlm_extension_base* clone() const {
        return new routingExtension(
            inputPortNumber, outputPortNumber);
    }

    void copy_from(const tlm_extension_base& ext) {
        const routingExtension& cpyFrom =
            static_cast<const routingExtension&>(ext);
        inputPortNumber = cpyFrom.getInputPortNumber();
        outputPortNumber = cpyFrom.getOutputPortNumber();
    }

    unsigned int getInputPortNumber() const {
        return inputPortNumber;
    }

    unsigned int getOutputPortNumber() const {
        return outputPortNumber;
    }
};
```

```
SC_MODULE(Interconnect)
{
    ...
    routingExtension* ext;

    ext = new routingExtension(
        inPort, outPort);

    trans.set_auto_extension(ext);
    ...
    routingExtension *ext = NULL;
    trans.get_extension(ext);
    outPort = ext->getOutputPortNumber();
    ...
};
```

Try code on github:
https://github.com/TUK-SCVP/SCVP.artifacts/blob/master/tlm_payload_extensions/

